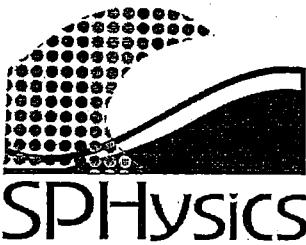


High-Performance Computing applied to Particle Hydrodynamics (SPH)



Benedict D. Rogers

School of Mechanical, Aerospace & Civil
Engineering, University of Manchester

@SPH_Manchester



UNIVERSITÀ
DEGLI STUDI
DI PARMA



TÉCNICO LISBOA



Universidade Vigo



SPH Course, University of Parma, 19-21 February 2019

Overview

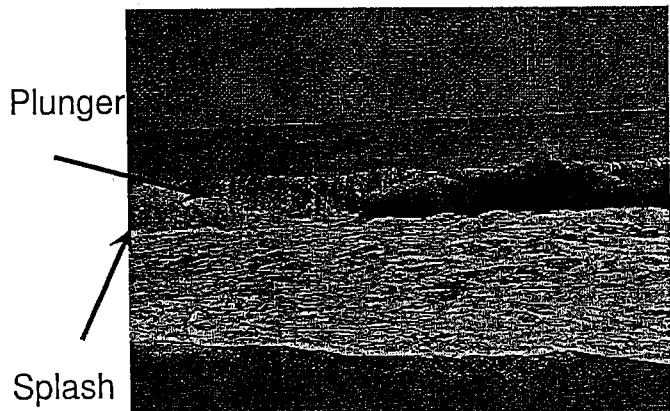
- Motivation, why SPH, why hardware acceleration is needed
- Basics of SPH algorithms: Neighbour Lists, Particle Interactions, System Update
- Options for hardware acceleration
- MPI, Domain Decomposition for SPH, Load Balancing and Optimisations
- The use of Graphics Processing Units (GPUs)
- Multi-Phase implementations
- Outlook



My Original Motivation for SPH

- Free-surface flows are rarely singly connected, e.g. beaches & wave energy devices

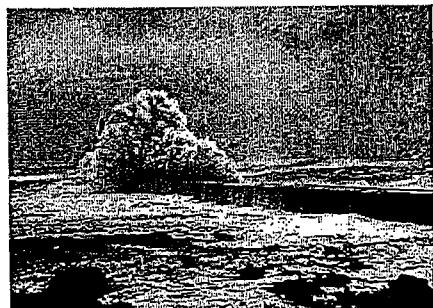
Breaking waves on beaches



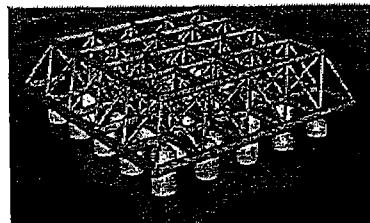
(Photo courtesy of F. Raichlen)

Very complex Multi-phase Multiscale problems

Overtopping:



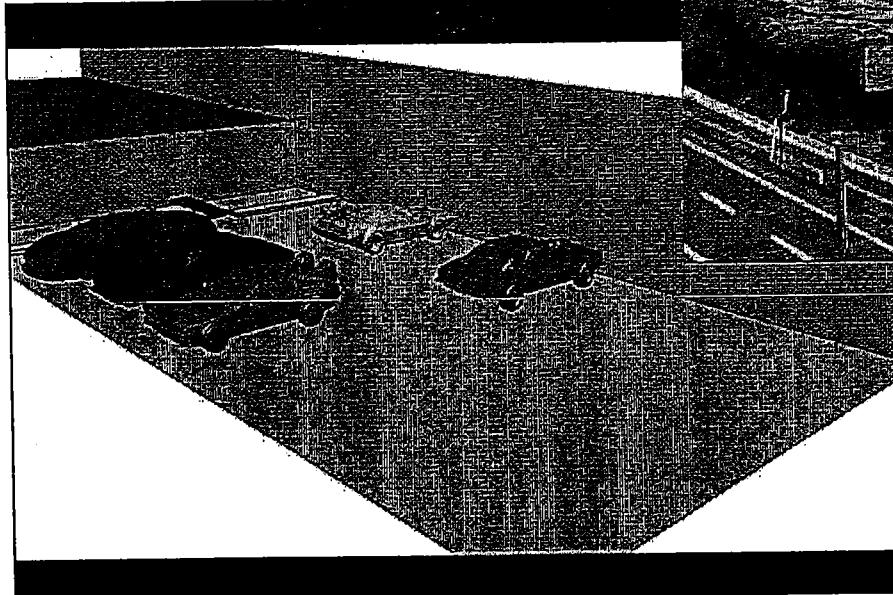
Wave Energy Devices:
Manchester Bobber



SPH: Has some distinct advantages in simulating these situations

Classical SPH Formulation Example

2011 Japanese Tsunami



Crespo et al. (2012)

Contents

- 1. Why does SPH require specialised computing?**
- 2. The 3 main steps of an SPH simulation**
- 3. Massive parallelisation with MPI**
- 4. Novel computing: GPUs**



Contents

- 1. Why does SPH require specialised computing?**
- 2. The 3 main steps of an SPH simulation**
- 3. Massive parallelisation with MPI**
- 4. Novel computing: GPUs**



Motivation for SPH simulations with large number of particles

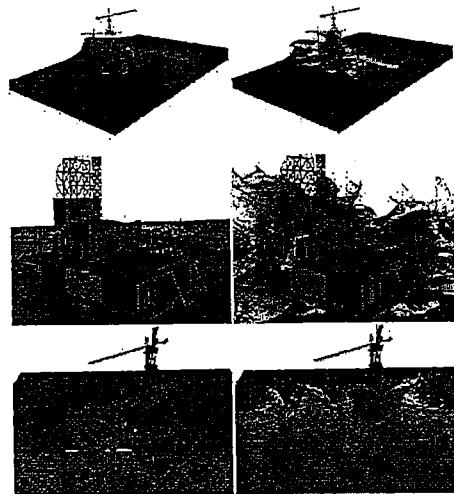
- Real engineering problems are 3-D, very large and/or act on multiple scales.

e.g. 1 billion particles of Jose on 64 GPUs with Weakly Compressible SPH

Dominguez et al. (2013)
Computer Physics Comm.

Very large particle numbers and high resolutions are required for high accuracy = **100+ million particles**

- An SPH Code should be able to scale on PetaFlops HPC platforms, with consideration of software development for Exascale.



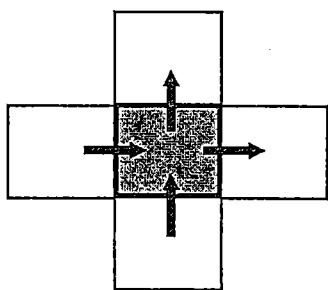
3-D applications require millions of particles

SPH was prohibitively expensive computationally

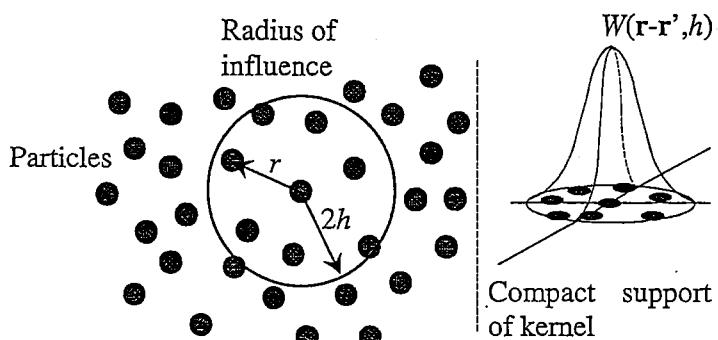
SPH simulations require acceleration, Why?

- The main problem is due to the interpolation procedure itself
- In Finite Volume, Finite Element & Finite Difference Schemes, the stencil around any cell usually contains only a small number of neighbouring cells! e.g. In 2-D FVM only 4 neighbouring cells

Finite Volume (FVM) Stencil



SPH Stencil



- In 2-D, each particle typically interacts with 20-50 particles
- In 3-D, each particle typically interacts with 100-400 particles

Contents

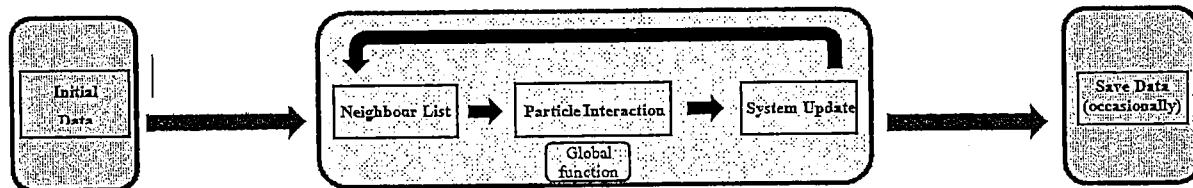
1. Why does SPH require specialised computing?
2. The 3 main steps of an SPH simulation
3. Massive parallelisation with MPI
4. Novel computing: GPUs

Let's remind ourselves of:

the 3 main steps of any SPH code

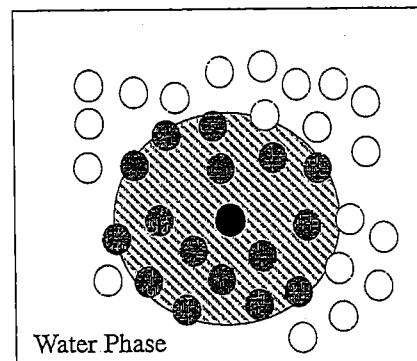
MANCHESTER
1824

Overview of an SPH code structure



We have 3 main steps:

1. Neighbour list (NL) creation
2. Particle Interaction (PI)
3. System Update (SU)



3 Main Steps of an SPH code

Step 1: Nearest neighbour searching

MANCHESTER
1824

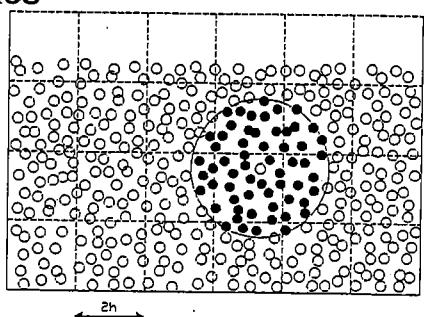
NEAREST NEIGHBOUR SEARCHING

In order to perform the summations required in SPH, we need to know which particles are near to each other

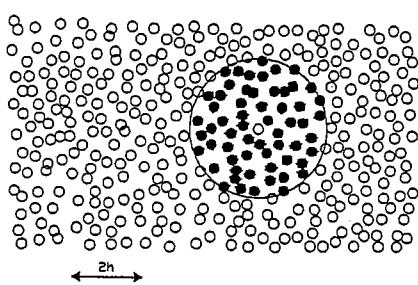
We cannot perform a summation over all the particles as this would take a prohibitive amount of time → hence the use of the compact support of the kernels

There are two general options for dividing up the computational domain efficiently:

(i) Divide up the domain into $2h$ boxes



(ii) Construct a Verlet list which is updated occasionally



NEAREST NEIGHBOUR SEARCHING

(i) Divide up the domain into $2h$ boxes

We split our domain into a set number of $2h$ boxes in each co-ordinate direction: ncx , ncy , ncz

We find in which $2h$ box a particle resides

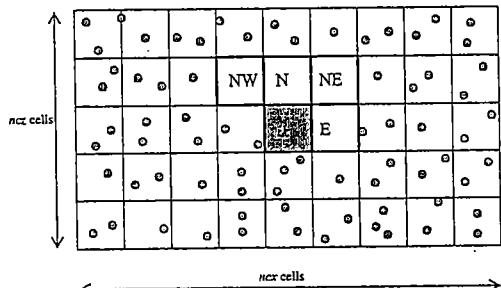
We make use of kernel symmetry (gradient asymmetry) to avoid repeating particle interactions

We sweep over the grid looking at particle interactions **between neighbour cells** where $r_{ij} < 2h$ (see SPHysics guide)

Pseudo-code:

$$\text{int}\left(\frac{x_i - x_0}{2h}\right) + 1 \Rightarrow \text{icell}$$

$$\text{int}\left(\frac{y_i - y_0}{2h}\right) + 1 \Rightarrow \text{jcell}$$



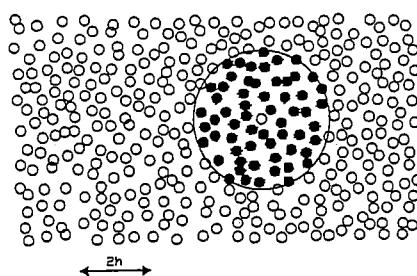
NEAREST NEIGHBOUR SEARCHING

(ii) Construct a Verlet list which is updated occasionally

We search our domain to create a list of neighbours for each particle that exist within $2h$ box or similar

Create a linked list and perform **sorting** algorithms to order particles for efficient sequential access

For more information see
Dominguez et al. (2010)



3 Main Steps of an SPH code

Step 2: FORCE COMPUTATION

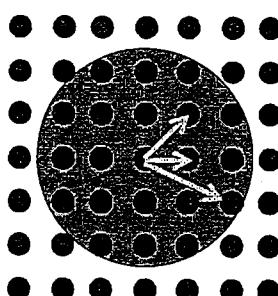
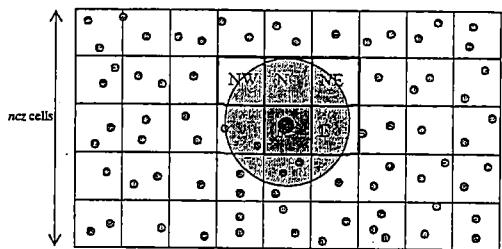
MANCHESTER
1824

FORCE COMPUTATION

This computes the acceleration terms in the governing equations and is the **MOST TIME CONSUMING PART**

$$\begin{aligned}\frac{d\rho_i}{dt} &= \sum_j m_j (\mathbf{v}_i - \mathbf{v}_j) \cdot \nabla_i W_{ij} \\ \frac{d\mathbf{v}_i}{dt} &= - \sum_j m_j \left(\frac{p_i}{\rho_i^2} + \frac{p_j}{\rho_j^2} \right) \nabla_i W_{ij} \\ &\quad + \sum_j m_j \frac{4\nu_o}{\rho_i + \rho_j} \frac{\mathbf{r}_{ij} \cdot \nabla_i W_{ij}}{r_{ij}^2 + 0.01h^2} (\mathbf{v}_i - \mathbf{v}_j) + \mathbf{F}_i\end{aligned}$$

We sweep through our $2h$ cells calculating our interactions for all the neighbours



$$\sum f_j \nabla_i W_{ij} = f_1 \nabla_i W_{i1} + f_2 \nabla_i W_{i2} + f_3 \nabla_i W_{i3} + \dots$$

3 Main Steps of an SPH code

Step 3: SYSTEM UPDATE: Time-stepping or Time integration

MANCHESTER
1824

SPH Time Integration

Let's see how we can march the simulation through time.

There is an enormous body of work on time stepping. We have a huge choice of schemes (some more sophisticated than others). As our particles (computation or nodal points) are moving we must use at least a scheme that is **2nd-order accurate in time**.

To begin with we rewrite all our governing equations in the following simplified form:

$$\left. \begin{array}{l} \frac{d\mathbf{v}_a}{dt} = \mathbf{F}_a \\ \frac{d\rho_a}{dt} = D_a \\ \frac{d\mathbf{r}_a}{dt} = \mathbf{v}_a \\ \frac{de_a}{dt} = E_a \end{array} \right\} \begin{array}{l} \text{Momentum} \\ \text{Continuity} \\ \text{Position} \\ \text{Energy} \end{array}$$

SPH Time Integration

2nd-order accurate Predictor-Corrector

Predictor step: This scheme predicts the evolution in time as

$$\mathbf{v}_a^{n+1/2} = \mathbf{v}_a^n + \frac{\Delta t}{2} \mathbf{F}_a^n \quad \rho_a^{n+1/2} = \rho_a^n + \frac{\Delta t}{2} D_a^n$$

$$\mathbf{r}_a^{n+1/2} = \mathbf{r}_a^n + \frac{\Delta t}{2} \mathbf{v}_a^n \quad e_a^{n+1/2} = e_a^n + \frac{\Delta t}{2} E_a^n$$

Corrector step: These values are then corrected using forces at half step

$$\mathbf{v}_a^{n+1/2} = \mathbf{v}_a^n + \frac{\Delta t}{2} \mathbf{F}_a^{n+1/2} \quad \rho_a^{n+1/2} = \rho_a^n + \frac{\Delta t}{2} D_a^{n+1/2}$$

$$\mathbf{r}_a^{n+1/2} = \mathbf{r}_a^n + \frac{\Delta t}{2} \mathbf{v}_a^{n+1/2} \quad e_a^{n+1/2} = e_a^n + \frac{\Delta t}{2} E_a^{n+1/2}$$

Finally, the values are calculated at the end of the time step

$$\mathbf{v}_a^{n+1} = 2\mathbf{v}_a^{n+1/2} - \mathbf{v}_a^n \quad \rho_a^{n+1} = 2\rho_a^{n+1/2} - \rho_a^n$$

$$\mathbf{r}_a^{n+1} = 2\mathbf{r}_a^{n+1/2} - \mathbf{r}_a^n \quad e_a^{n+1} = 2e_a^{n+1/2} - e_a^n$$

Contents

1. Why does SPH require specialised computing?
2. The 3 main steps of an SPH simulation
3. Massive parallelisation with MPI
4. Novel computing: GPUs

Let's look at how the 2 main hardware acceleration options work

First, the CHOICE of HARDWARE determines the options of parallelisation technique ...

MANCHESTER
1824

Hardware Acceleration: the options

- (i) Using parallel (supercomputer) machines with lots of cores (individual CPUs) and divide the work over them

Qu: What's the difference between a parallel machine and a supercomputer?

- (ii) FPGAs – Field Programmable Gate Arrays: well used in astrophysics simulations, but expensive and not portable

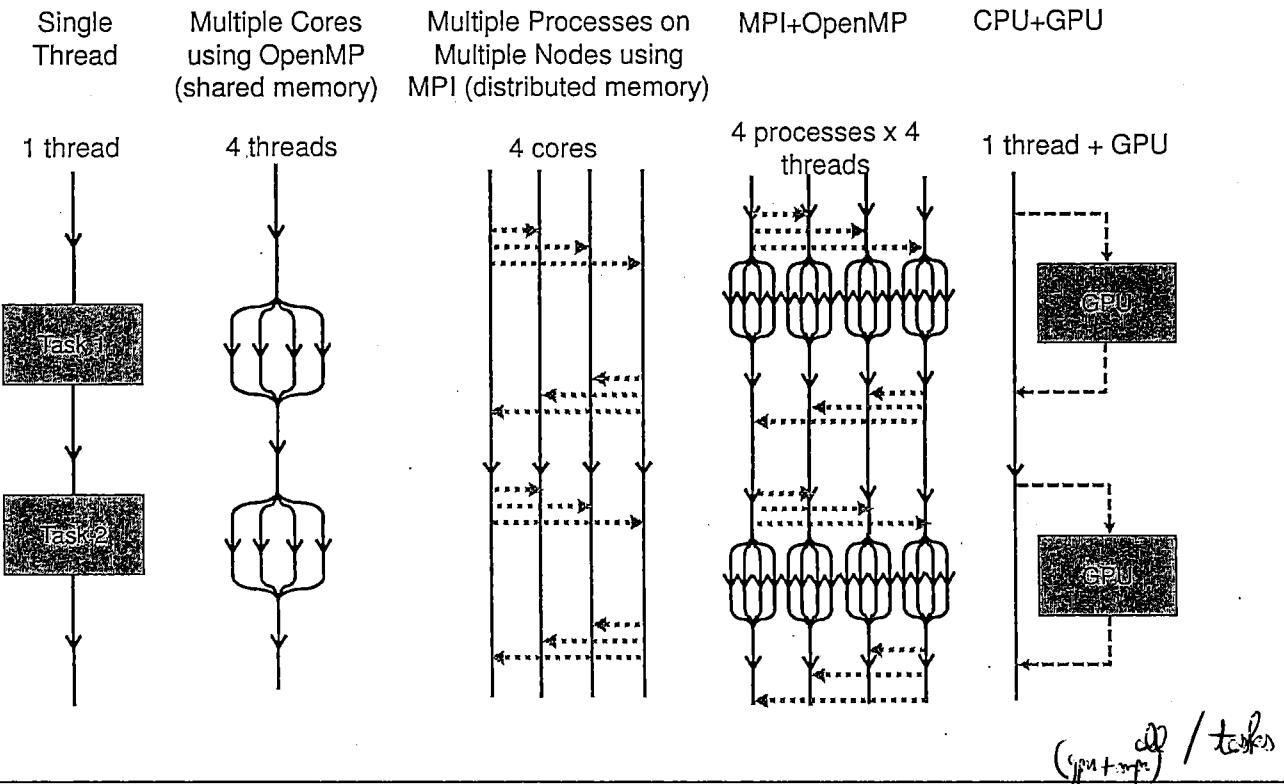


- (iii) GPUs – Graphics Processing Units: the hot topic of scientific computing



Simulation Acceleration: the options

Depending on hardware 5 basic **PROGRAMMING** options exist:



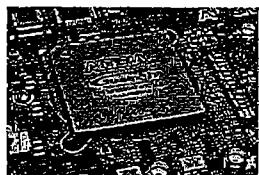
Hardware Acceleration: the options

- (i) Using parallel (**supercomputer**) machines with lots of cores (individual CPUs) and divide the work over them

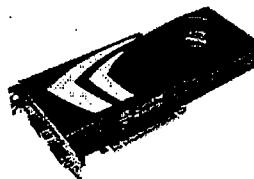
Qu: What's the difference between a parallel machine and a supercomputer?

→ interconnect speed = super

- (ii) **FPGAs** – Field Programmable Gate Arrays: well used in astrophysics simulations, but expensive and not portable



- (iii) **GPUs** – Graphics Processing Units: the **hot topic** of scientific computing



Hardware Acceleration for SPH

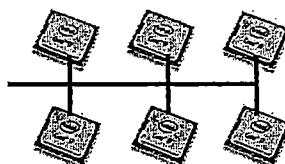
Option 1: Parallel computing using Distributed Memory with MPI

Parallel Algorithm Development

How to speed up a computation? 1 processor is not enough!



So, use lots of processors!



Spread the computational load between many processors

This requires communication

In the 1990s, the Message Passing Interface (MPI) emerged as the standard approach to achieve this communication:

Data is transferred between different processors using MPI commands:

e.g. to send data MPI_SEND
to receive data MPI_RECV

Parallel Algorithm Development

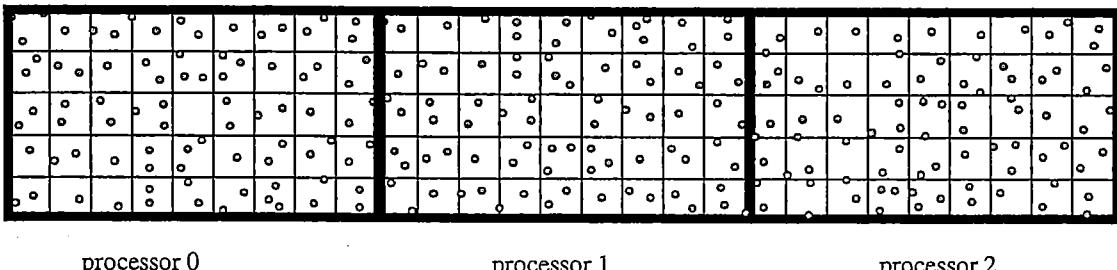
There are 3 key steps for massive parallelisation:

1. **Domain Decomposition** to distribute the load of the computation among different processors
2. **Load Balancing** so processors are not left doing nothing
3. **Optimisation:** Reducing cache access time

Key Step 1: Domain Decomposition

Parallel Algorithm Development

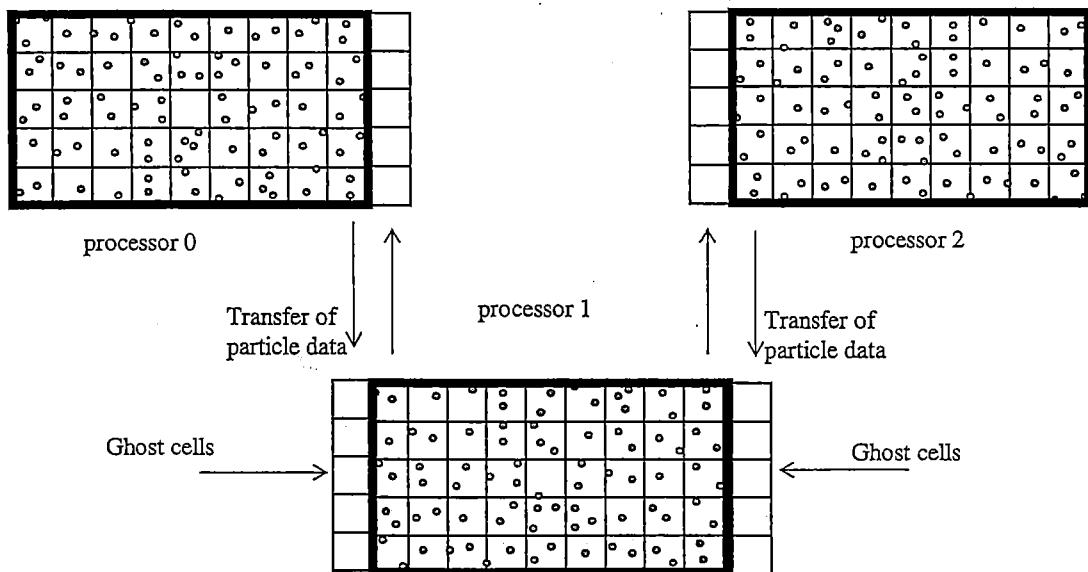
Connectivity of the particles is achieved using the common technique of splitting the domain into 2^h square boxes



Subdivision of the domain across processors

Parallel Algorithm Development

Communication is then necessary



Data is exchanged between each processor using ghost cells

This is ideal for simple rectangular domains that are completely full

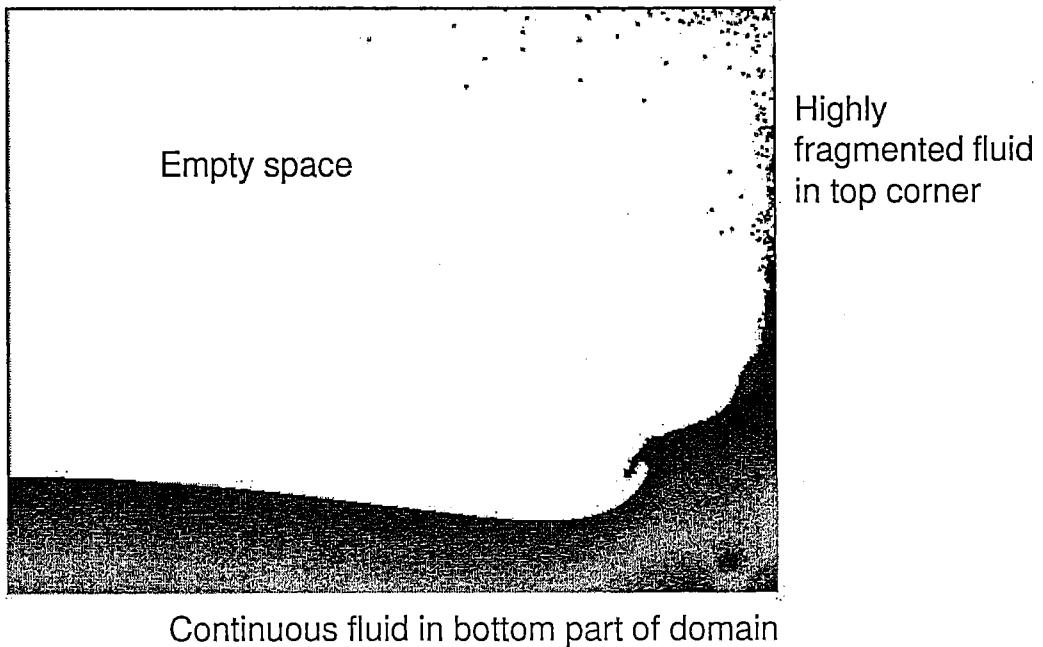
But it's very limited for SPH!

**More sophisticated
Domain Decomposition**

Guo et al. (2018) CPC

Why is a more sophisticated domain decomposition needed?

- Many of our applications will have highly fragmented fluid phases



How do we use our computing resources efficiently where particles are located?

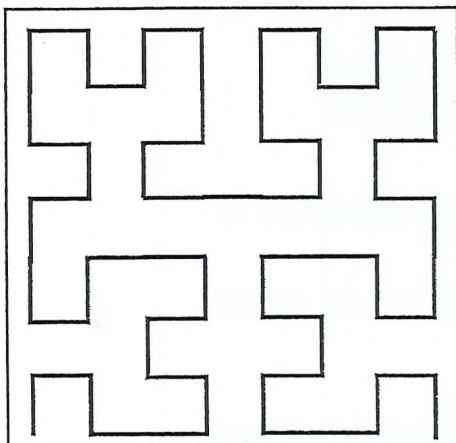
- {
 - Octree
 - ORB
 - Space Filling Curves

Generating Space Filling Curve

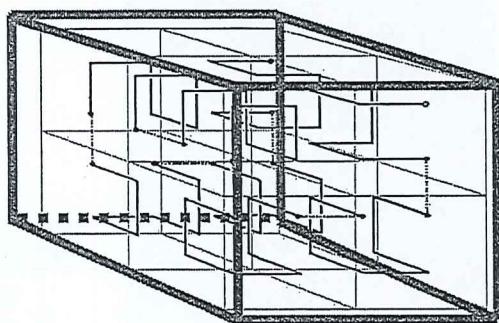
How do you map a space efficiently?

SPACE FILLING CURVES: Hilbert Space Filling Curve

2-D

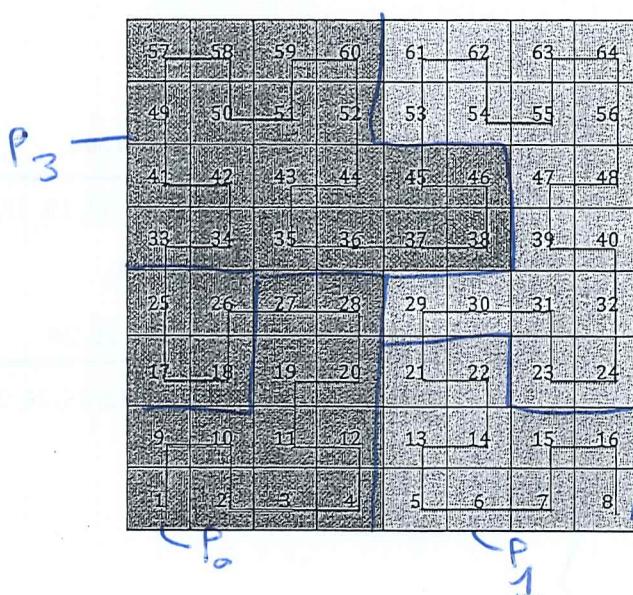


3-D



How to use a Space Filling Curve for Domain Decomposition

1. Cover the domain using the Space Filling Curve
2. Use a Domain Decomposition algorithm (e.g. Zoltan) to split the domain into subdomains



Partitions: P_0 =blue,
 P_1 =green
 P_2 =yellow
 P_3 =pink

Next Question:

How do particles located on P_0 interact with particles on other Partitions?

Halo Creation

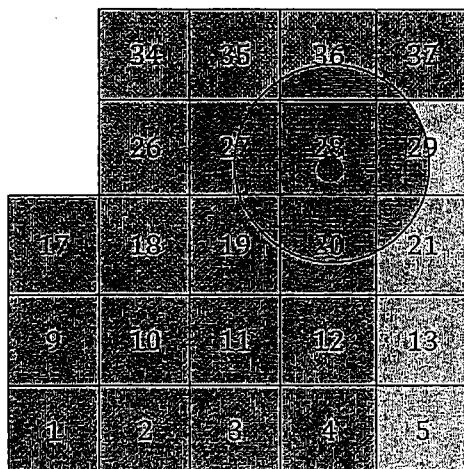
- Consider Particles located in Partition P_0 . They need the particle information in the cells located on other partitions

Partitions: P_0 =blue,

P_1 =green

P_2 =yellow

P_3 =pink



Consider a particle located in cell 28.

It will have neighbours in

other cells:

{ 19, 20, 21, 27, 28, 29, 35,
36, 37 }

Hence we create a HALO

Halo Creation: MPI_SEND

- Zoltan for Domain Decomposition & halo
- Hilbert Space Filling Curve

Partitions: P_0 =blue,

P_1 =green

P_2 =yellow

P_3 =pink

Halo Send and Receive Plan for cells:

P0:SEND(to send to P3): 9, 10, 11, 19, 27, 28

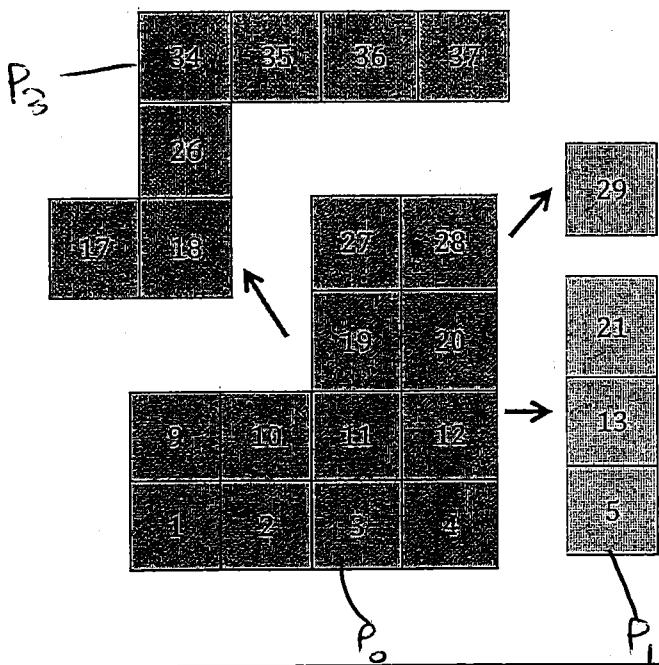
P0:SEND(to send to P2): 20, 28

P0:SEND(to send to P1): 4, 12, 20, 28

P0:RECV(from P3): 17, 18, 26, 32, 35, 36, 37

P0:RECV(from P2): 29

P0:RECV(from P1): 5, 13, 21



Halo Creation: MPI_RECV

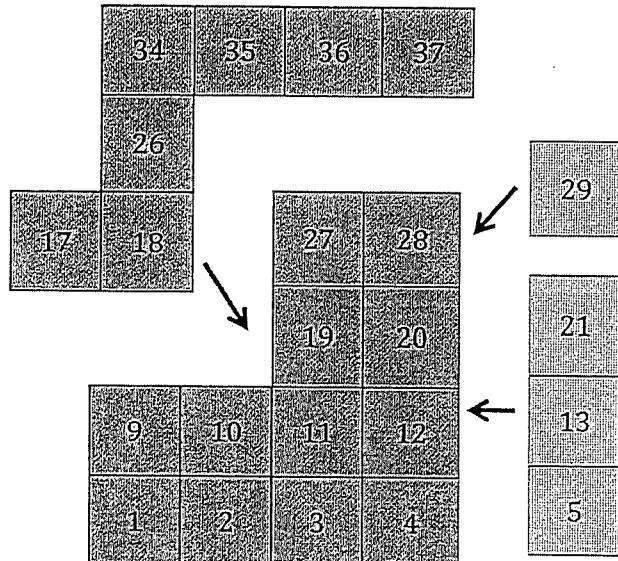
- Zoltan for Domain Decomposition & halo
- Hilbert Space Filling Curve

Partitions: P_0 =blue,

P_1 =green

P_2 =yellow

P_3 =pink



Halo Send and Receive Plan for cells:

P0:SEND(to send to P3): 9, 10, 11, 19, 27, 28

P0:SEND(to send to P2): 20, 28

P0:SEND(to send to P1): 4, 12, 20, 28

P0:RECV(from P3): 17, 18, 26, 32, 35, 36, 37

P0:RECV(from P2): 29

P0:RECV(from P1): 5, 13, 21

Key Step 2: Load Balancing

Load Balancing

There are 2 key principles of load balancing for any parallel code:

- Make Each Partition do the same amount of work
- Reduce Communication

This requires

- Load Balancing Strategy
- Sophisticated Communication (overlapping communication, etc.)

The aim is to give **SCALABILITY** over 1000s of cores

Load Balancing: Zoltan

ZOLTAN is a library for massively parallel systems:

<http://www.cs.sandia.gov/Zoltan/>

Each cell, i , has a weight is computed as:
where NP = number of particles

$$wgts(i) = \frac{NP(i)}{NP_{total}}$$

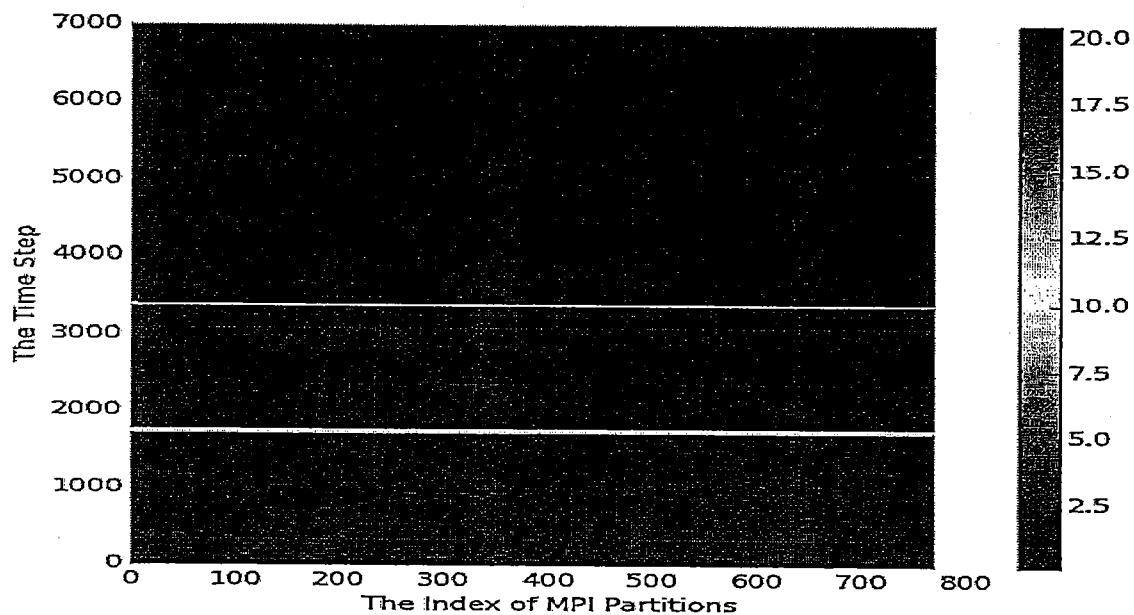
Load on each processor = sum of cell weights = $load = \sum_{i=1}^{N_k} wgts(i)$

Imbalance for each processor = $imbalance = \frac{load_{max}}{load_{average}} = \frac{\max(load)}{NP_{total}/K}$

If Imbalance > 20% tolerance, Use ZOLTAN to partition the cells along the
HSFC = Hilbert Space Filling Curve
(See Guo et al. 2018 CPC).

Load Balancing: Heat Maps

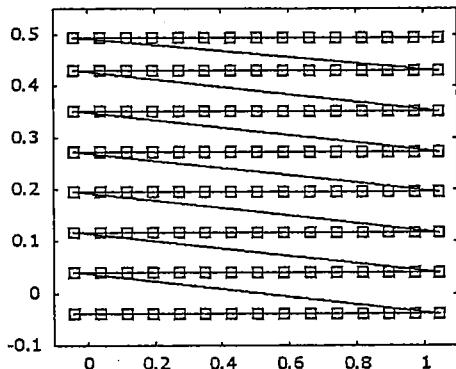
The easiest way to check that your code is well balanced is using HEAT MAPS:



Key Step 3: Reducing Cache Access Time

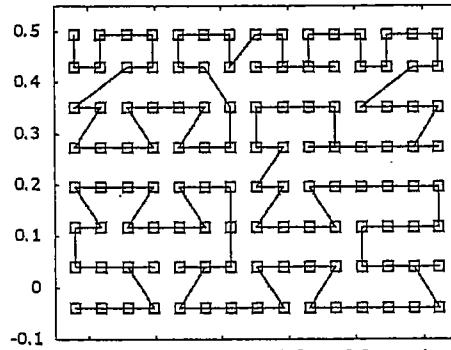
Particles-Cells reordering

The *inherited* cell linked list of particles would sweep through the domain in a sequential order:



The natural order of cells

To improve *cache* performance and avoid overchecking of neighbours use Hilbert Space Filling Curve [HSFC] to reorder the list of particles for neighbours only:



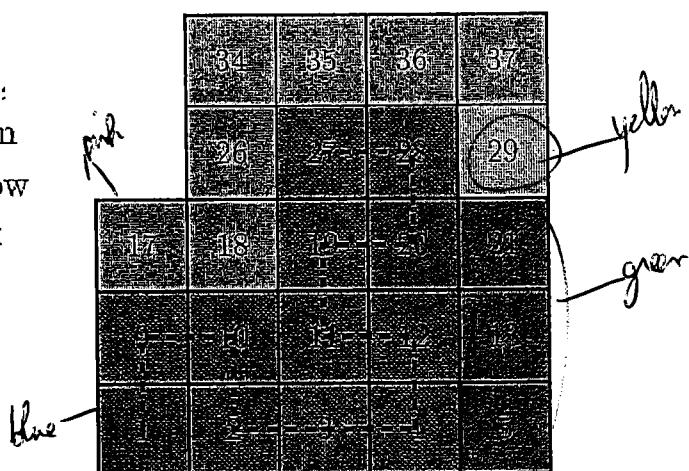
The HSFC order of cells

(see Oger et al. 2010)

Reordering Cells on a Partition

- The order of the particles is important to maximise the time it takes to retrieve the values of particle properties from memory (CACHE ACCESS)

Partitions: P_0 =blue,
 P_1 =green
 P_2 =yellow
 P_3 =pink

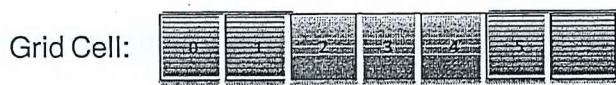


Cell's Natural Order: 1 2 3 4/9/10 11 12/19 20 27 28 = \rightarrow
Cells HSFC Order: 1 9 10 2 3 4 12 11 19 20 28 27

4 cells to 3, 12 but not 9!

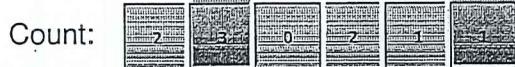
Cell Linked List: preconditioned dynamic vector and linked list

Particle arrangement
within the cells



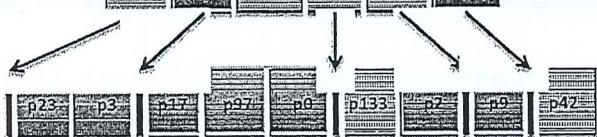
↓
Apply
pre-conditioning

New list of particles stored for
faster access



Dominguez et al. (2011)

Storage:



Reduces Memory Footprint: crucial for using massive HPC systems

General Performance Considerations for SPH Parallelization

We don't just want a highly parallelised SPH scheme, we also need to solve the Pressure Poisson Equation over 10-30,000 cores!

The challenges can be summarised as:

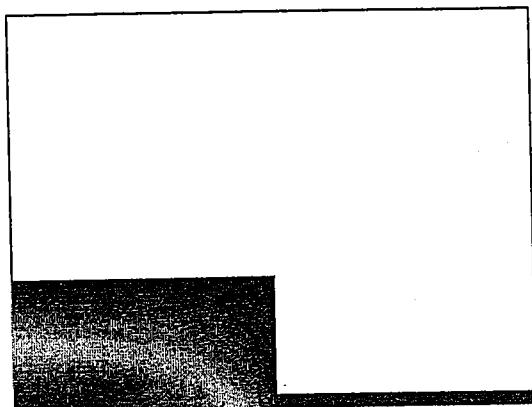
- Efficient neighbour list searching
- Depends upon a good assignment of particles to processors;
 - Grouping physically close particles within a single processor
 - Reduces inter-processor communication
 - Re-partitioning of the particles to maintain geometric locality of objects within processors
- Low load balancing costs for non-uniformly distributed particles arising from solving a complex, highly nonlinear and distorted flow and adaptive SPH.

+ Strong & Weak Scaling

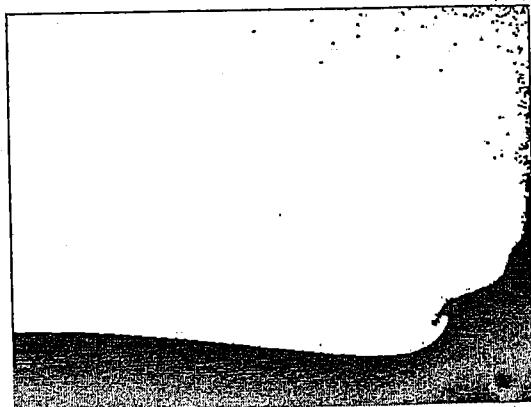
Demonstration Case: 2-D Wet-Bed Dam Break

What does this look like when we run it for a wet-bed dam break?

An example of a violent nonlinear flow requiring highly adaptive domain decomposition across empty space:



$t = 1$



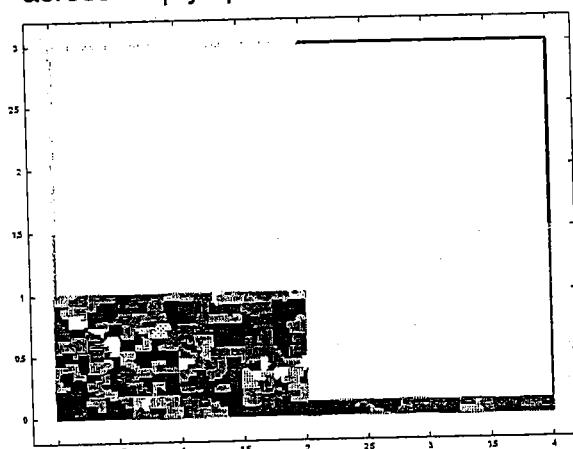
$t = 5720\Delta t$

Colours denote pressure

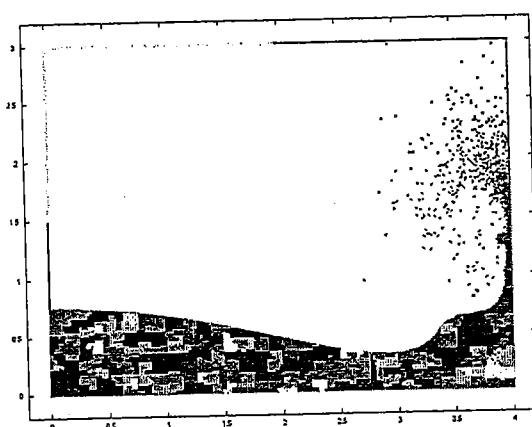
Dynamic Domain Decomposition with HSFC

What does this look like when we run it for a wet-bed dam break?

An example of a violent nonlinear flow requiring highly adaptive domain decomposition across empty space:



256 partitions at $t = 1$



256 partitions at $t = 5720\Delta t$

Colours denote partitions

What about Incompressible SPH (ISPH)?

MANCHESTER
1824

Incompressible SPH method

- ISPH_DF – A Divergence-free Projection method in SPH

$$\mathbf{r}_i^* = \mathbf{r}_i^n + \delta t \mathbf{u}_i^n$$

Particles are advected to \mathbf{r}_i^*

$$\mathbf{u}_i^* = \mathbf{u}_i^n + (\nu \nabla^2 \mathbf{u}_i^n + \mathbf{F}_i^n) \delta t$$

Calculate \mathbf{u}_i^* without pressure gradient.

$$\nabla \cdot \left(\frac{1}{\rho} \nabla p^{n+1} \right)_i = \frac{1}{\delta t} \nabla \cdot \mathbf{u}_i^*$$

Solve pressure Poisson equation to get pressure P_i^{n+1} .

$$\mathbf{u}_i^{n+1} = \mathbf{u}_i^* - \frac{\delta t}{\rho} \nabla p_i^{n+1}$$

Calculate velocity \mathbf{u}_i^{n+1} .

$$\mathbf{r}_i^{n+1} = \mathbf{r}_i^n + \delta t \left(\frac{\mathbf{u}_i^{n+1} + \mathbf{u}_i^n}{2} \right)$$

Particles' positions are centered in time.

$$\delta \mathbf{r}_s = -\mathcal{D} \left(\frac{\partial C}{\partial s} \mathbf{s} + \alpha \left(\frac{\partial C}{\partial n} - \beta \right) \mathbf{n} \right)$$

Shift particles to maintain stability.

*ISPH (Cummins & Rudman 1999)

Parallel ISPH Solver features

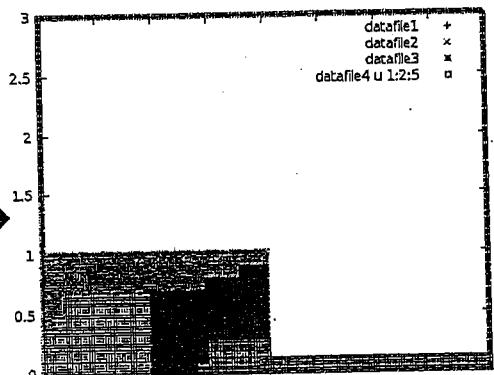
In addition to the ISPH formulation

- ISPH utilizes Zoltan for spatial **Domain Decomposition** and **Dynamic Load Balancing** for evenly distributed use of computational resources.
- To solve Poisson Pressure Equation

$$\nabla \cdot \left(\frac{1}{\rho} \nabla p^{n+1} \right)_i = \frac{1}{\delta t} \nabla \cdot \mathbf{u}_i^* \rightarrow -\frac{1}{\rho} \sum_j V_j (p_j - p_i) \nabla W_{ij}$$

ISPH use PETSc (Portable, Extensible Toolkit for Scientific Computation).

- Fortran 90, MPI.



General Performance Considerations for **ISPH** Parallelization

We don't just want a highly parallelised SPH scheme, we also need to solve the Pressure Poisson Equation over 10-30,000 cores!

The challenges can be summarised as:

- Efficient neighbour list searching
- Depends upon a good assignment of particles to processors;
 - Grouping physically close particles within a single processor
 - Reduces inter-processor communication
 - Re-partitioning of the particles to maintain geometric locality of objects within processors
- Low load balancing costs for non-uniformly distributed particles arising from solving a complex, highly nonlinear and distorted flow and adaptive SPH.
- Efficient and scalable pressure Poisson solver

Solving Pressure Poisson Equation(PPE) with PETSc

$$\nabla \cdot \left(\frac{1}{\rho} \nabla p^{n+1} \right)_i = \frac{1}{\delta t} \nabla \cdot \mathbf{u}_i^*$$



After cell & particle reordering Let's solve the Pressure Poisson Equation = matrix eq

$$\mathbf{A}\mathbf{X} = \mathbf{b} \quad -----$$

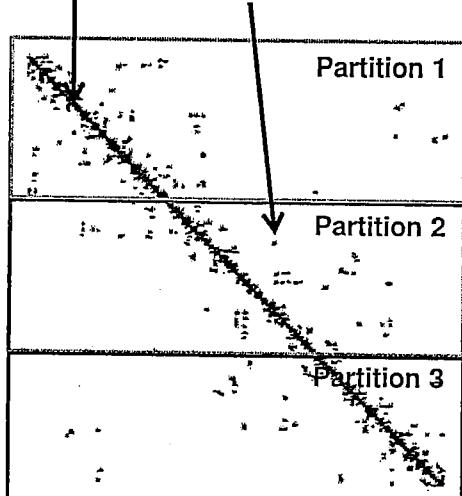
$$a_{ii}P_i + \sum \alpha_i a_{ij}P_j = \alpha_i b_i$$

ISPH uses PETSc for the sparse linear solver.

- the non-zeros of each row in the matrix represent each particle's neighbour particles
- The PETSc algorithm partitions the matrix by continuous rows, which requires a renumbering of particles so that insertions of values to global matrix become local operations without incurring extra communications

$$\mathbf{AX} = \mathbf{b} \Leftrightarrow \mathbf{PAPTPX} = \mathbf{Pb}$$

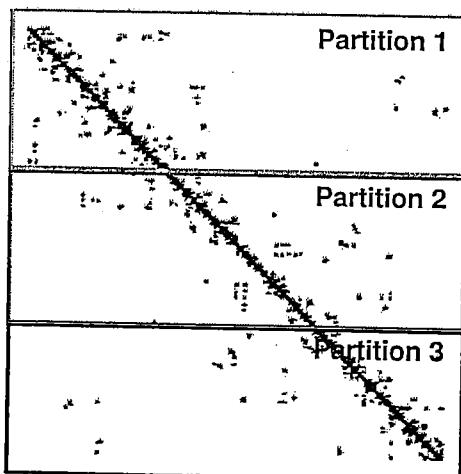
- benefit of reducing the bandwidth of coefficients matrix



Sparse Pattern of PPE Matrix

Problem with convergence of the PETSc matrix solver:

- Many solvers can be slow to converge
- Previously used **Jacobi**(preconditioner) +Bi-CGSTAB for single core but this is prohibitively slow.
- With a large number of particles (100+ million) there can be **no convergence** at all.
- For example, with **Jacobi+BiCGSTAB** won't converge, often requires 10,000+ iterations to converge every Δt .



Sparse Pattern of PPE Matrix

Key Question for ISPH for 100+ million particles over 10,000 cores:

How to address convergence for 100+ million particles for sparse PPE matrix? Need good/efficient preconditioner

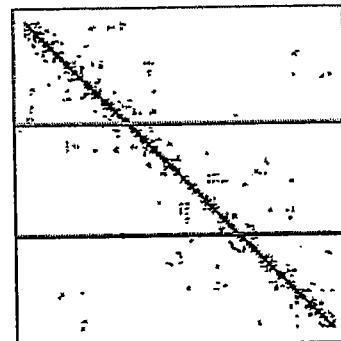
SUPERFAST preconditioning for convergence of the PETSc matrix solver:

SOLUTION?

How to address converge for 100+ million particles?

Multigrid preconditioner:

HYPRE BoomerAMG



Sparse Pattern of PPE Matrix

BoomerAMG has 2 phases:

1. setup(PCSETUP)

selection of coarse grids

creation of the interpolation operators

the representation of the fine grid matrix operator on each coarse grid

2. solve(PCAPPLY)

matrix-vector multiply

+ the smoothing operator

ISPH PPE Solver Performance Analysis: strong scaling

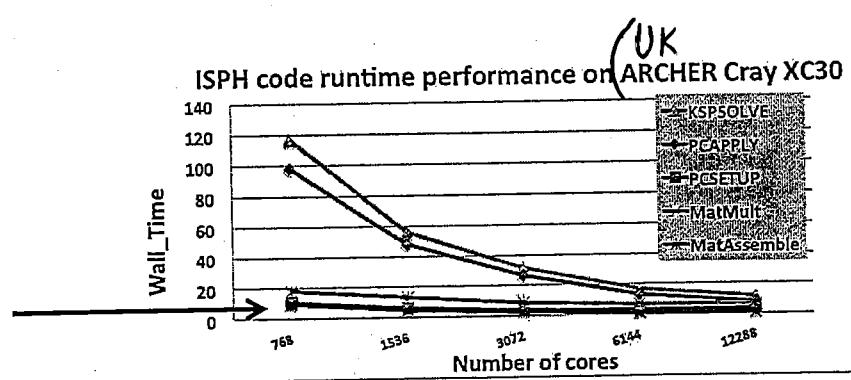
Strong scaling = fixed number 100million particles but vary number of cores

But with more cores → more communication and this normally hurts your computation runtime ☺ a bit challenging

Question 1: What is the pre-conditioning's cost in PETSc for ISPH PPE?

Wall clock time for pre-conditioning operations (PCSETUP) is very small for all # cores,

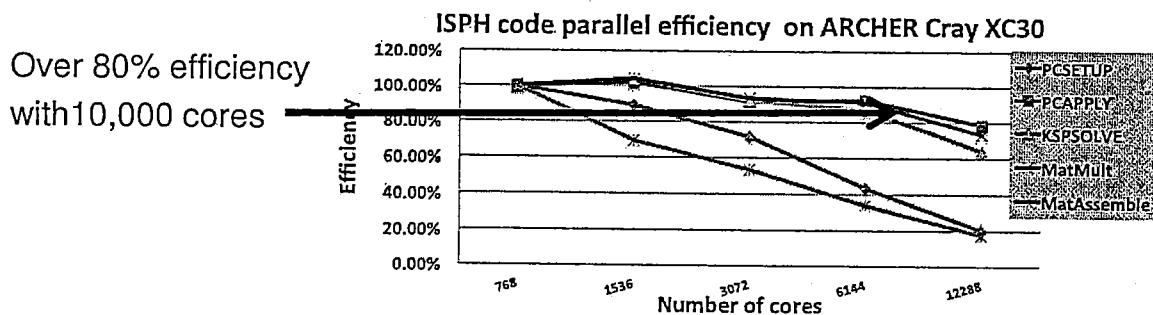
PCAPPLY(smoothen and SpMV) has the same cost with solver phase (KSPSOLVE).



ISPH PPE Solver Performance Analysis: strong scaling

Strong scaling = fixed number 100million particles but vary number of cores

Question 2: How does multigrid preconditioner improve ISPH performance?



Number of iterations for convergence changed from

10,000+ → 12 iterations per timestep!!

Answer: quite promising for 100+million particles
... but may not be enough for ExaSCale

Some Key Results from MASSIVELY PARALLEL SPH CODES

For SPH:

- 40 million – EDF using Spartacus on a Blue Gene
- 120 million – EPFL using a Blue Gene
- 200 million – ECN using SPH-flow in France
- 1 billion – UK Astrophysics Consortium
- 1 billion – DualSPHysics Consortium
- 1 billion – SWIFT astrophysics
- 1.2 billion – Karlsruhe, Germany

Impressive, BUT these simulations needed massive HPC supercomputers. This raises important questions:

- What are the limitations?
- Are these accessible to all potential users/developers/researchers of SPH?

Limits of an SPH Parallel Codes

- What is the **limit** of what we can simulate on supercomputing clusters?
- Another way of asking that question is how many particles can we realistically use in a simulation before we get hit by prohibitive communication and just the sheer number of processors that we need to support **millions-billions** of particles
- Thankfully, there is a new solution!!!!

Contents

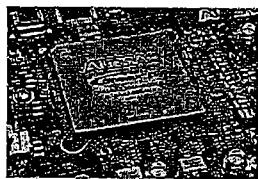
1. Why does SPH require specialised computing?
2. The 3 main steps of an SPH simulation
3. Massive parallelisation with MPI
4. Novel computing: GPUs

Hardware Acceleration: the options

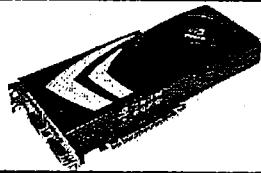
- (i) Using parallel (**supercomputer**) machines with lots of cores (individual CPUs) and divide the work over them

Qu: What's the difference between a parallel machine and a supercomputer?

- (ii) FPGAs – Field Programmable Gate Arrays: well used in astrophysics simulations, but expensive and not portable



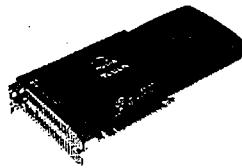
- (iii) GPUs – Graphics Processing Units: the **hot topic** of scientific computing



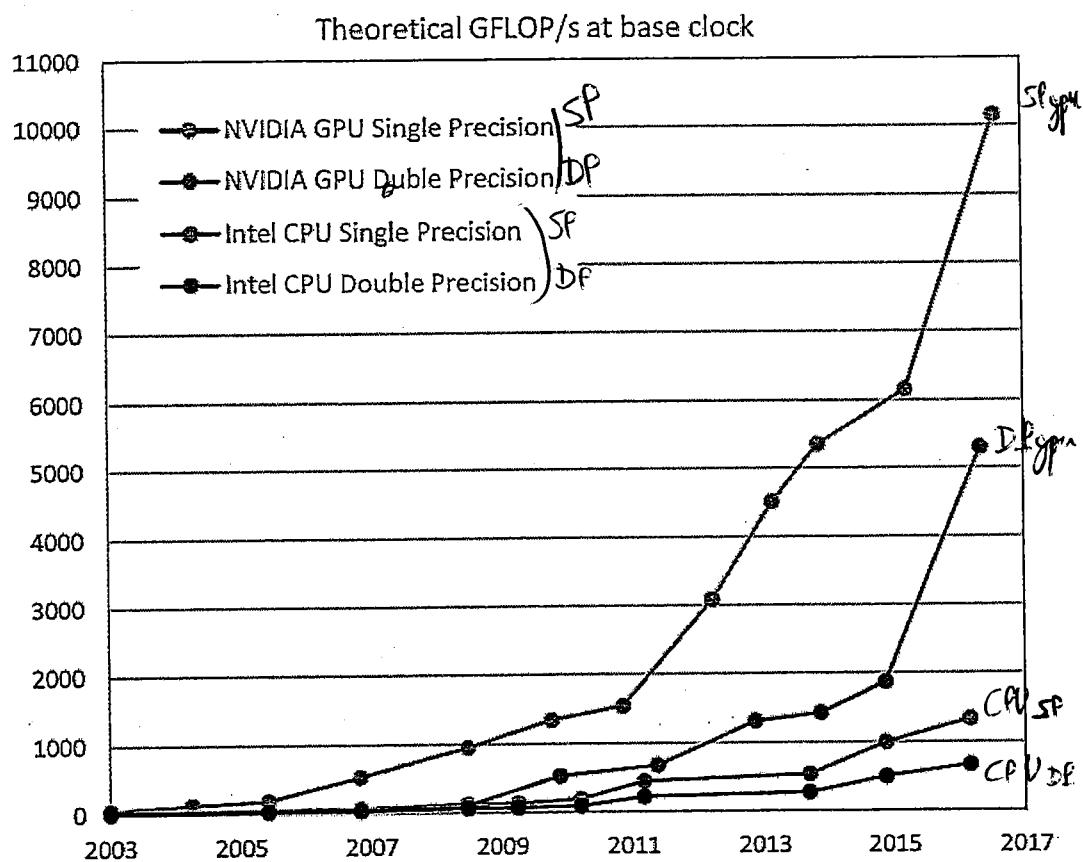
Hardware Acceleration for SPH

Option 2: GPUs

What is a GPU?



- **GPU = Graphics Processing Unit**
 - They are used in all your computers for the graphics
- **GPUs developed independently of the CPU**
(CPU=Central Processing Unit used for normal computation)
 - This means the architecture is very different
 - In order to play Games in real time, they process a lot of information very quickly
 - Hence, they consist of many **parallel threads**
- **nVidia GPUs have an dedicated language - CUDA**
 - CUDA = Compute Unified Device Architecture
 - Make sure you buy a a **CUDA-enabled GPU!!!**
 - Programming is far more difficult as it is a lower level language than C/C++ & Fortran
 - You must know how to use the architecture
- **nVidia produce specific GPUs for Scientific Computing: FERMI cards**



Proposals to accelerate SPH

How to accelerate your SPH code with less than

£250

Intel® Core™ i7 940 at
2.93GHz with 4 cores



Multi-core
OpenMP

Speedup: 4x

£350

GTX 680 at 1GHz
with 1536 cores



GPU
CUDA

Speedup: 60-90x

£15,000

2 Intel Xeon E5620 + 4 GTX680



Multi-GPU
CUDA + MPI

Speedup: 100-500??

TOP SUPERCOMPUTERS IN THE WORLD Nov 2018

<http://www.top500.org>

R_{max} and R_{peak} values are in TFlops. For more details about other fields, check the TOP500 description.

TOP 10 Cities November 2018

Rank	Site	System	Cores	[TFlop/s]	[TFlop/s]	(kW)
1°	Sunway Taihulight (China)	93.0 petaflops/s (consumption: 15371 KW)				(CPUs)
3°	Titan (USA)	17.5 petaflops/s (consumption: 8209 KW)				(with GPUs)

Energy Efficient GPU co-processors are
now a key component in HPC

Part of computing Emerging Technology
2019 Conference:
<http://emit.tech>

High Performance
Computing
Japan

M1; Intel Xeon Phi 7250 88C 1.4GHz, Intel
Omni-Path
Fujitsu

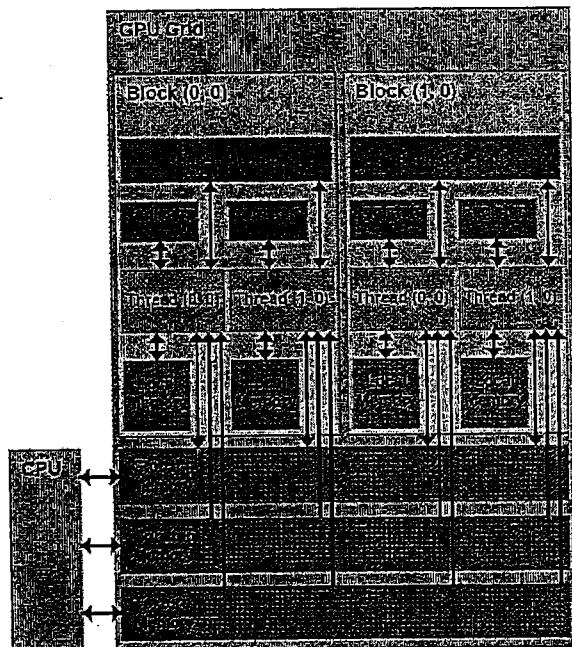
09-11/04/2019
UK

Hoffler

Factors affecting Performance Issues on GPUs

- Several memory spaces with different access modes and performance (shared memory, global memory, texture memory, ..)
- Branch divergence (avoid if statements)
- Non-balanced workload (not all streaming multi-processors finish simultaneously)
- Streaming Multi-Processor Occupancy
- No coalescent memory access (this is how data is grouped together on the GPU, 32, 64 or 128 bytes)

Memory Hierarchy



Inside each SM, there are **REGISTERS** which are fast and hold data but are small in number

The GPU generally consists of Streaming Multiprocessors (SMs) with access to different types of memory:

Global Memory: largest & slowest

Constant Memory: Read-only resides in Global Memory

Texture Memory: in Global memory optimised for 2-D spatial data locality

Local Memory: used only when registers have all been used

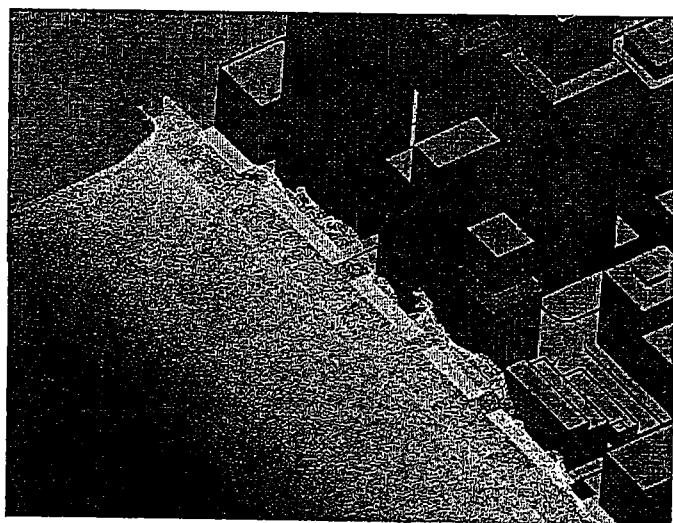
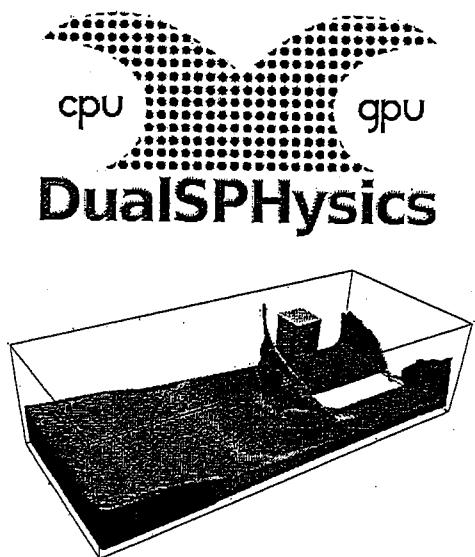
Shared Memory: is FAST as it is located in each Streaming Multiprocessor (SM)

**So, programming a GPU is not
easy but also not so difficult**

**There are now several GPU
codes ...**

MANCHESTER
1824

**DualSPHysics,
new GPU computing on SPH models**



A.J.C. Crespo, J.M. Dominguez, A. Barreiro and M. Gómez-Gesteira
EPHYSLAB, Universidade de Vigo, SPAIN

B. D. Rogers and D. Valdez-Balderas
MACE, The University of Manchester, U.K.

Universidade de Vigo

MANCHESTER
1824

SPH Momentum Equation

$$\left(\frac{d\mathbf{u}}{dt} \right)_i = - \sum_j m_j \left(\frac{p_j}{\rho_j^2} + \frac{p_i}{\rho_i^2} \right) \nabla_i W_{ij} + \mathbf{g}$$

Set of fluid particles

Smoothing domain
neighbor particles j

particle of interest i

SPH Momentum Equation

SUMMATION SOLUTION

CPU / Fortran

GPU / CUDA

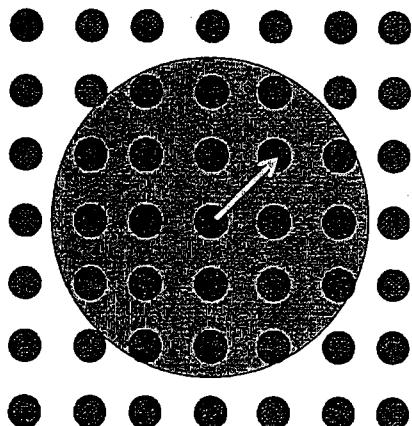
$$i \rightarrow j \quad \boxed{\mathbf{a}_i} \quad \left(\frac{d\mathbf{u}}{dt} \right)_i = - \sum_j m_j \left(\frac{p_j}{\rho_j^2} + \frac{p_i}{\rho_i^2} \right) \nabla_i W_{ij} + \mathbf{g}$$

SPH Momentum Equation

SUMMATION SOLUTION

CPU / Fortran

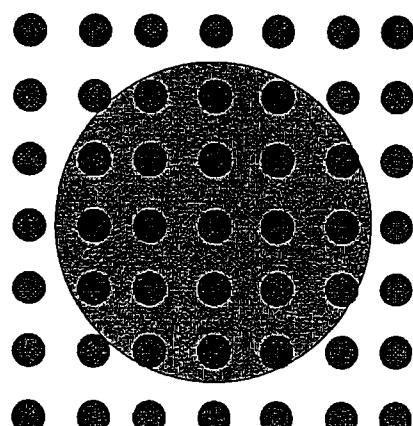
GPU / CUDA



$$i \rightarrow j = 1$$

$$a_{ij=1}$$

$$a_i = a_{i1} + \dots$$

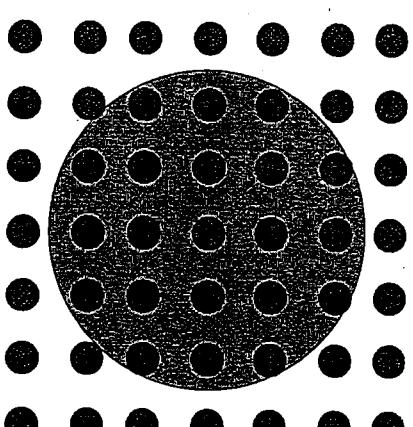
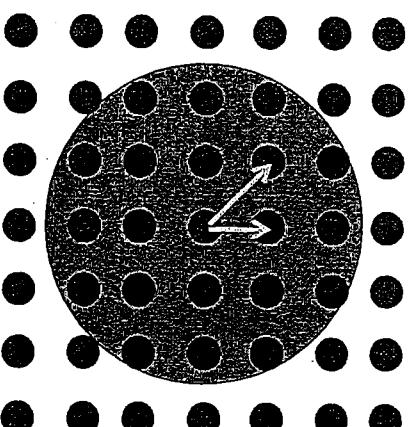


SPH Momentum Equation

SUMMATION SOLUTION

CPU / Fortran

GPU / CUDA



$$i \rightarrow j = 1$$

$$i \rightarrow j = 2$$

$$a_{ij=2}$$

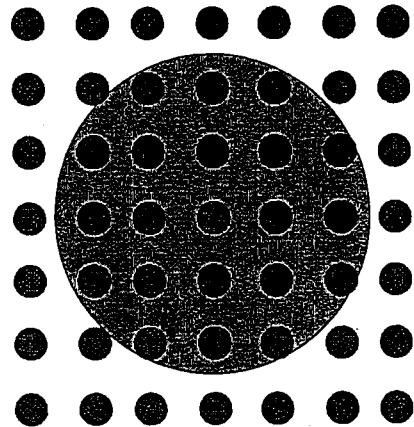
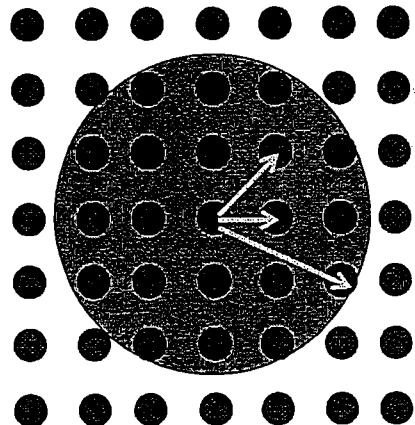
$$a_i = a_{i1} + a_{i2} + \dots$$

SPH Momentum Equation

SUMMATION SOLUTION

CPU / Fortran

GPU / CUDA



$$i \rightarrow j = 1$$

$$i \rightarrow j = 2$$

$$i \rightarrow j = 3$$

$a_{ij=3}$

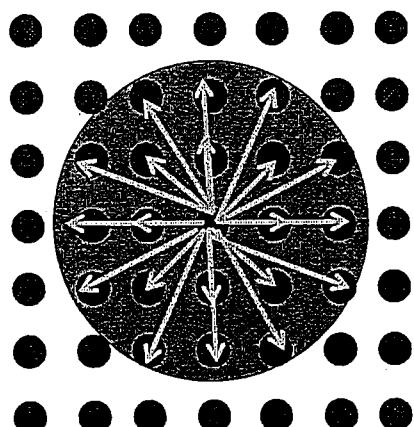
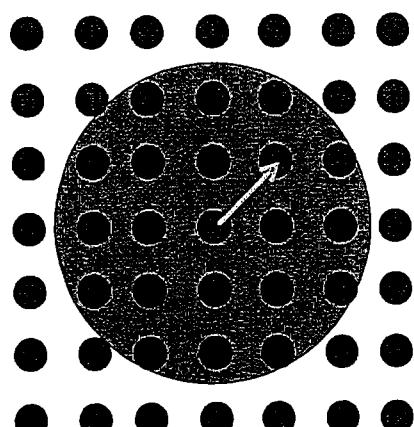
$$a_i = a_{i1} + a_{i2} + a_{i3} + \dots$$

SPH Momentum Equation

SUMMATION SOLUTION

CPU / Fortran

GPU / CUDA

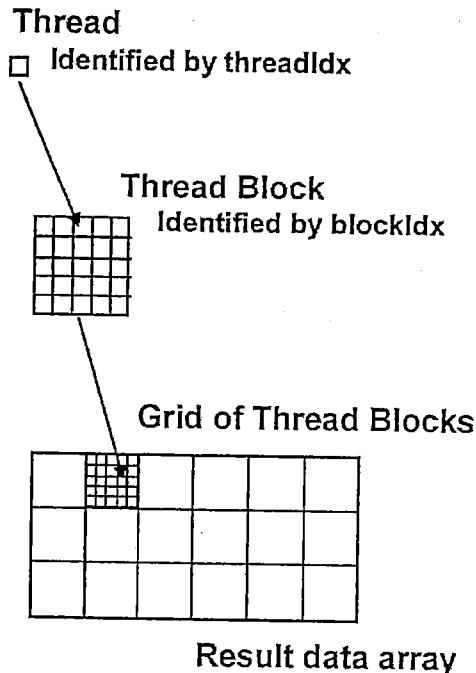


$$i \rightarrow j = 1 \quad a_{ij=1}$$

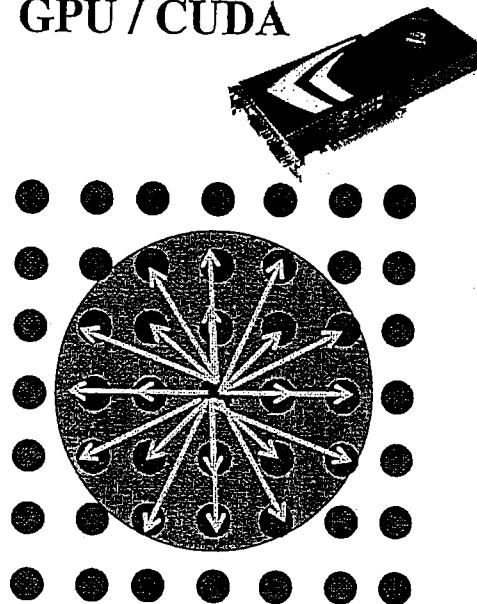
$$a_i = a_{i1} + \dots$$

$$i \rightarrow j \quad a_{ij}$$

$$a_i = \sum a_{ij}$$



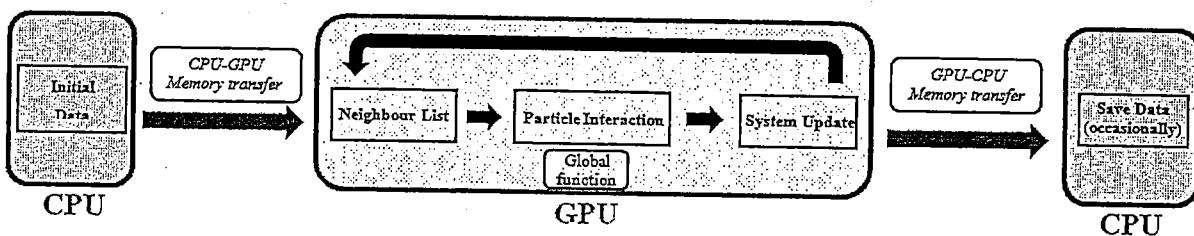
SUMMATION SOLUTION GPU / CUDA



Speedups of
20-100

$$\left(\frac{du}{dt} \right)_i = \sum_j m_j \left(\frac{p_j}{\rho_j^2} + \frac{p_i}{\rho_i^2} \right) \nabla_i W_{ij} + g \quad i \rightarrow j$$

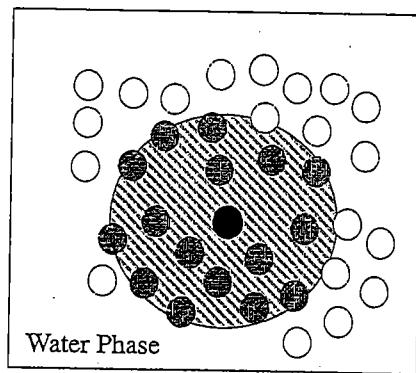
Overview of GPU code structure



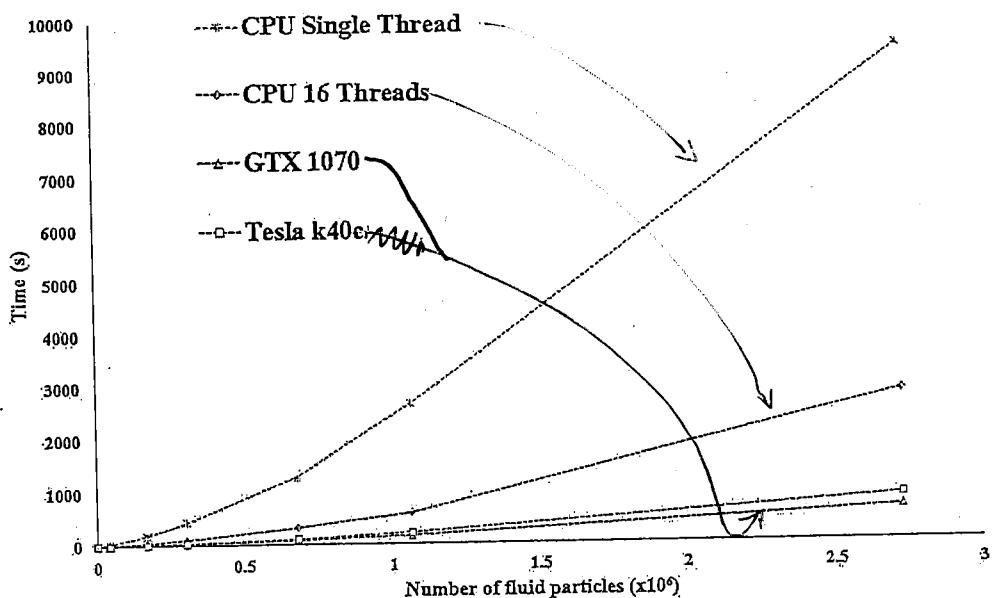
We have 3 main steps:

1. Neighbour list (NL) creation
2. Particle Interaction (PI)
3. System Update (SU)

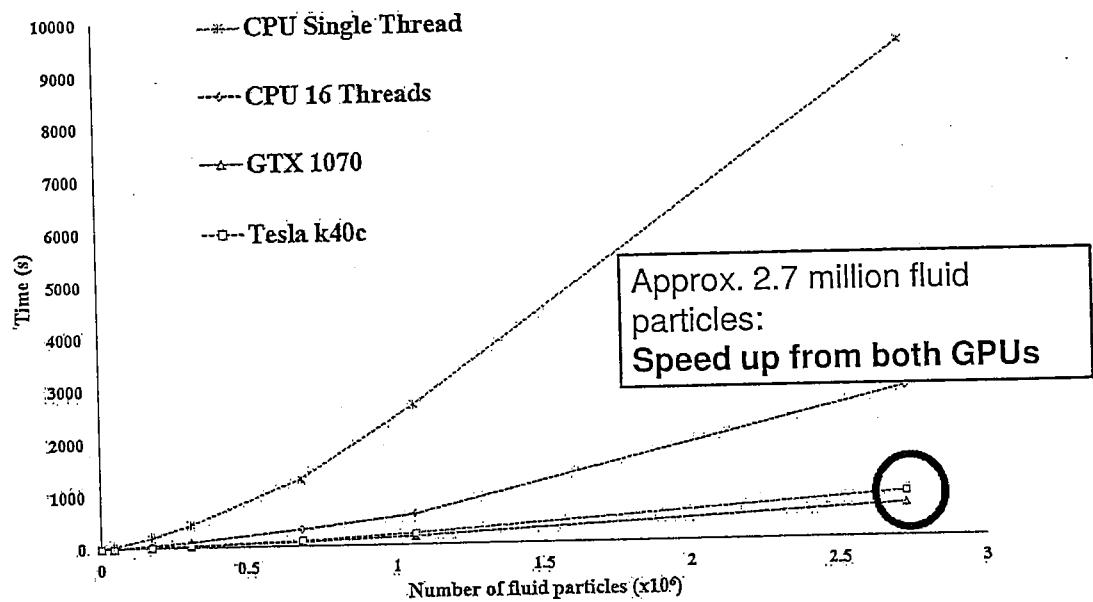
All within 1 single global function



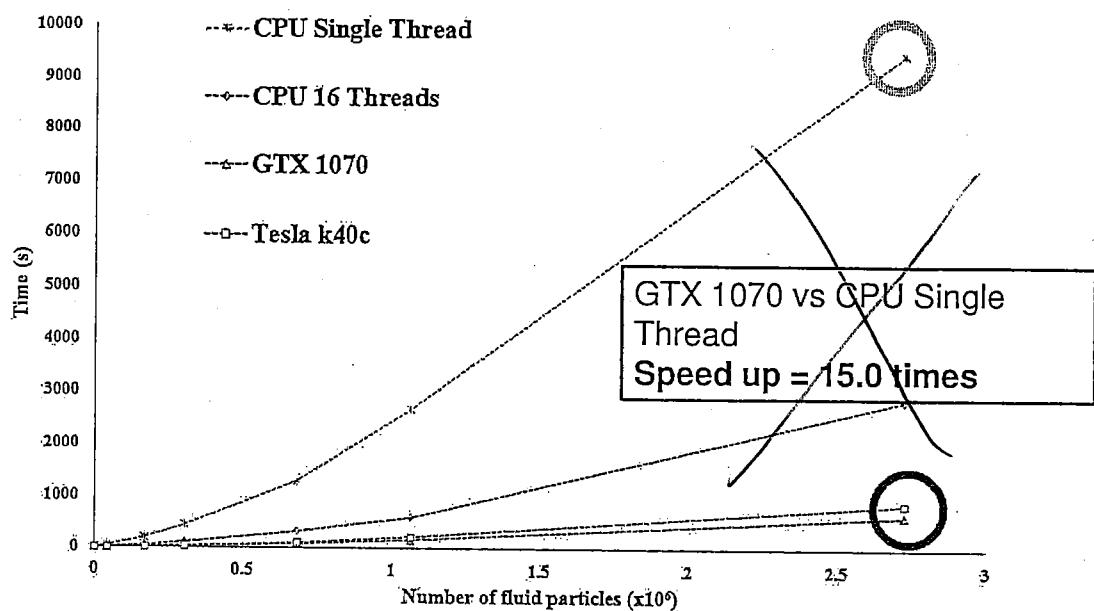
2D Dambreak runtime comparisons



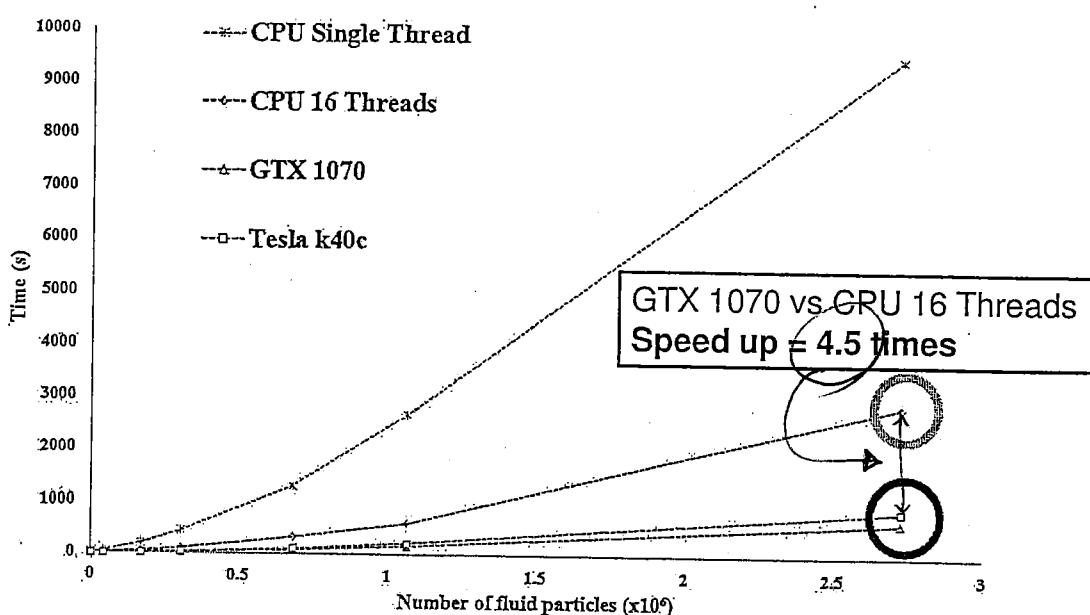
ISPH on a GPU (Chow et al. 2018)



ISPH on a GPU (Chow et al. 2018)



ISPH on a GPU (Chow et al. 2018)



ISPH on a GPU (Chow et al. 2018)

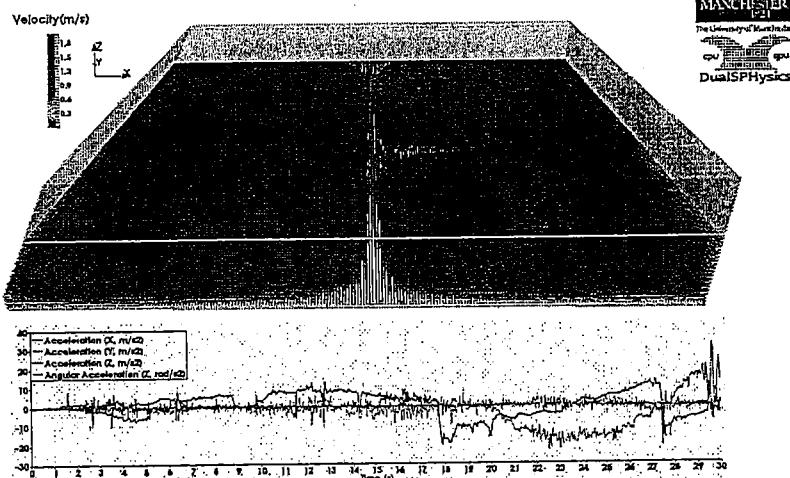
Fuel-tank sloshing with Leading Motorsport Company (F1)

Real engineering problems are now accessible

Only allowed to show
highly simplified
geometry (NDA)

Accelerations are up
to $5g$ \approx fighter jet

Comparisons with
in-tank footage were
close.



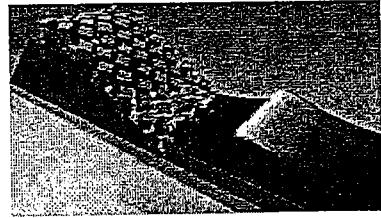
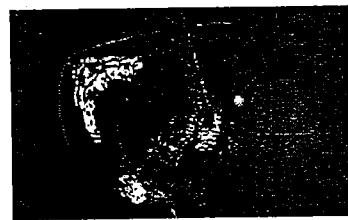
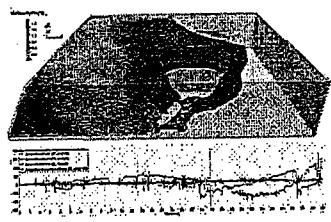
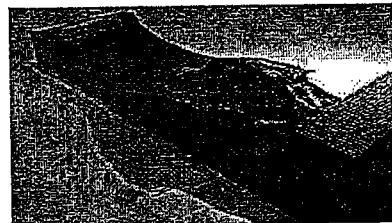
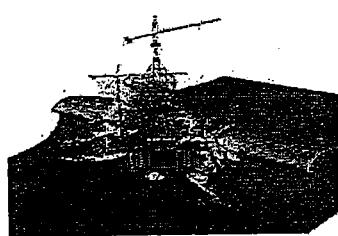
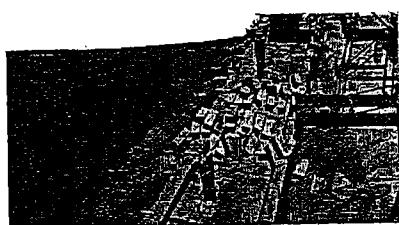
- Last race became engine stopped (even with Fuel) \Rightarrow know fuel was far from pump in the tank!



Longshaw & Rogers (2015), Advances Engineering Software

Funded by Knowledge Transfer Account (KTA), now the IAA

DualSPHysics Project:
runs on multi-core CPU or GPU



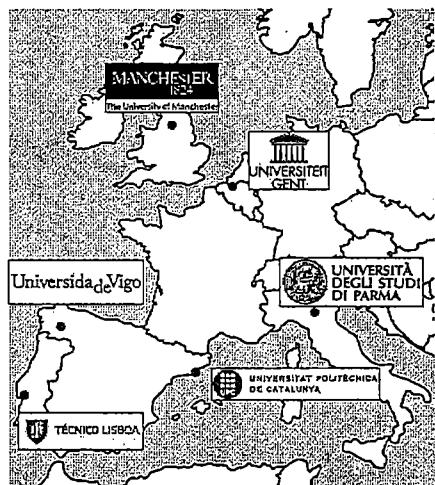
2nd DualSPHysics Users Workshop

University of Manchester · 13-14 September 2016

5th Users Workshop
March 2020, Barcelona, Spain

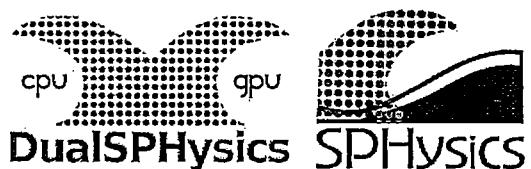
DualSPHysics Project:

- University of Manchester
- University of Vigo (Spain)
- University of Parma (Italy)
- University of Lisbon (Portugal)
- University of Ghent (Belgium)



Websites

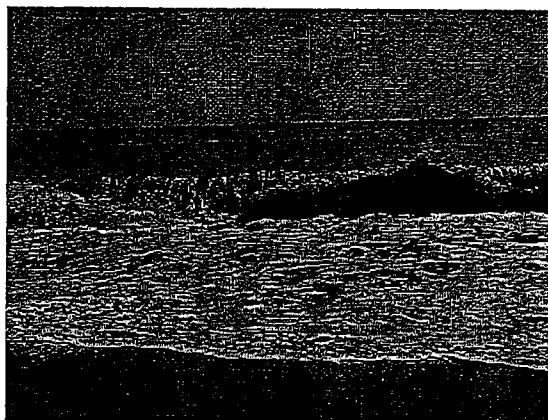
- Free open-source SPHysics code:
<http://www.sphysics.org>
<http://www.dual.sphysics.org>



Downloaded 15,000+ times: Open-source plug & play SPH code for free-surface flow

Okay, BUT

Everything there was mono-phase



(Photo courtesy of F. Raichlen)

MANCHESTER
1824

SPH free-surface Applications

Multi-Phase Modelling: WATER & AIR

Mokos *et al.* (2015, 2017) on GPUs

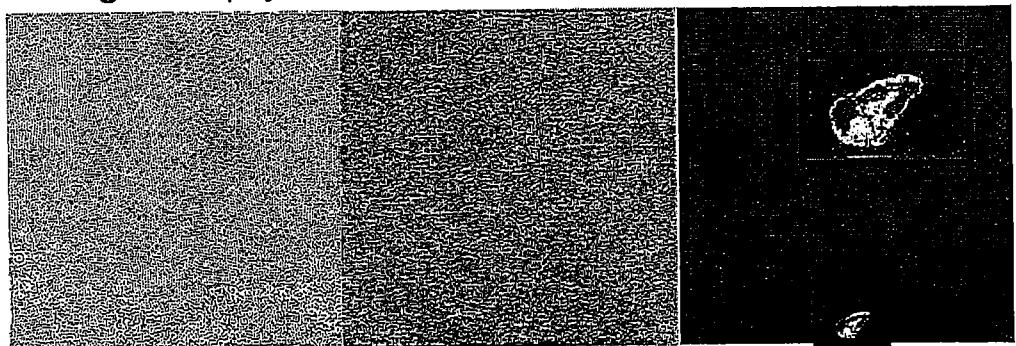
MANCHESTER
1824

Challenges of a multi-phase GPU code

- **Challenge 1:** Interaction of a gas and a liquid phase
 - Large density ratio (~ 1000)
 - Large pressure gradients in the interface
 - Treatment of the gas phase
- **Challenge 2:** Computational treatment
 - SPH is computationally expensive
 - Increased number of particles due to the second phase
- **Challenge 3:** Unphysical Voids Created at Higher Resolutions

Use appropriate formulation
(Yildiz' talk)

Use a GPU



Challenges of a multi-phase GPU code

Issue: Voids appear only in high resolutions

Solution: Fickian-based approach by Lind et al. (2012)

- Numerical treatment based on Fick's law:

$$\delta \mathbf{r}_s = -D \nabla C_i$$

- Shifting dependent on particle concentration: $\nabla C_i = \sum_j C_{ij} \frac{m_j}{\rho_j} \nabla W_{ij}$

- Diffusion based on particle velocity (Skillen et al. 2013): $D = -A_s h \|u\|_i \Delta t$

- Free-surface correction term:

- Used only for the liquid phase

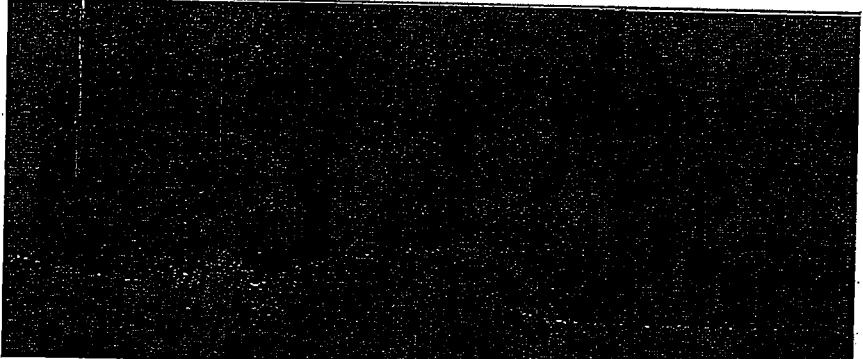
$$\delta \mathbf{r}_s = -D \left(\frac{\partial C_i}{\partial s} \mathbf{s} + \frac{\partial C_i}{\partial b} \mathbf{b} + \alpha \left(\frac{\partial C_i}{\partial n} - \beta \right) \mathbf{n} \right)$$

Wet Dam Break

- Original Result



- Particle Shifting



Demands of a multi-phase GPU code

Distinguish particles belonging to different phases

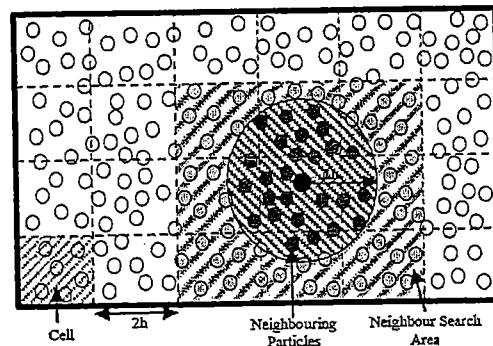
- Load different initial data for each phase
- ID-system recognising phase of each particle (use of *mkvalues*)

Optimise the multi-phase model

- Different equations used for each phase

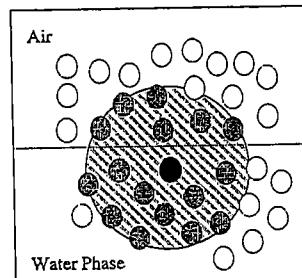
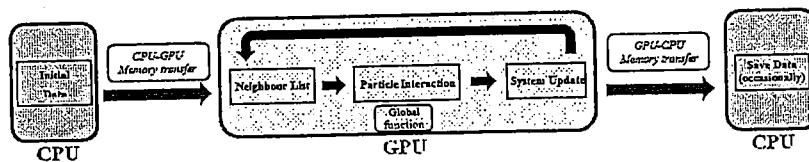
Maintain the existing structure of DualSPHysics

- Integration with other capabilities of the code, such as motion
- Maintain the efficient cell-linked-list structure
- Figure 1: Example of the Cell-linked List using 2h×2h 2h Cell Neighbour Search Area 2h Neighbouring Particles



Demands of a multi-phase GPU code

Could we use the original code structure?

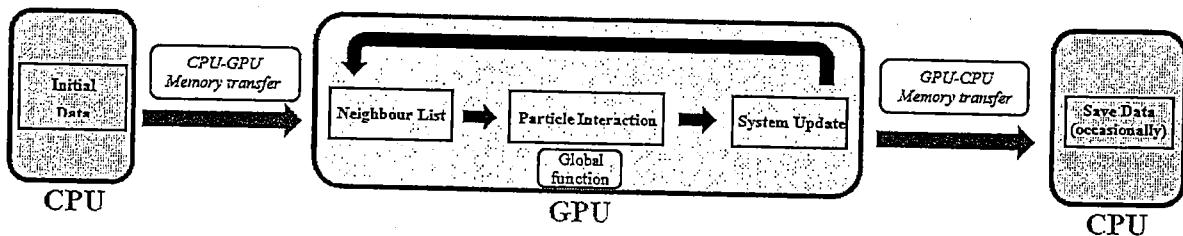


Well, in effect we have 4 choices:

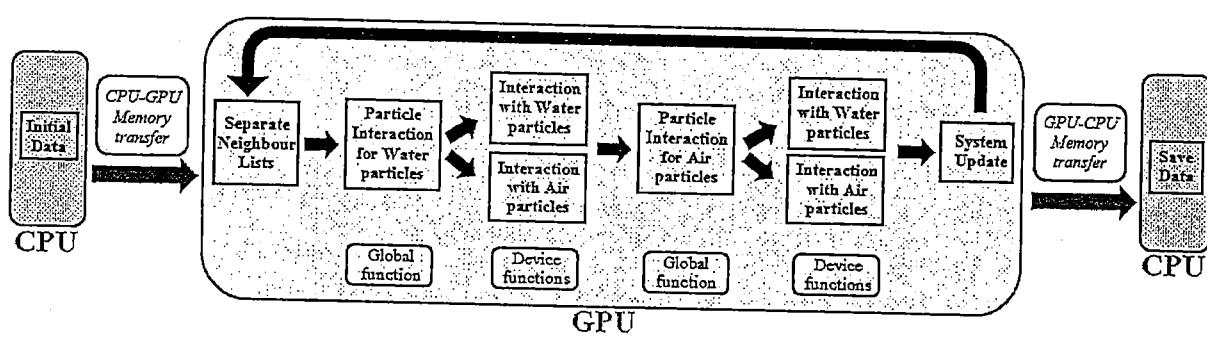
1. Use *if*-statements to distinguish each phase
2. Use binary multipliers (0 or 1 for each phase)
3. Separate neighbour lists + separate global & device functions
4. Separate neighbour lists + Intermediate CPU-GPU function
(avoids repeated CPU function calls)

Multi-phase GPU code: separate neighbour lists 2

Original Mono-Phase structure:

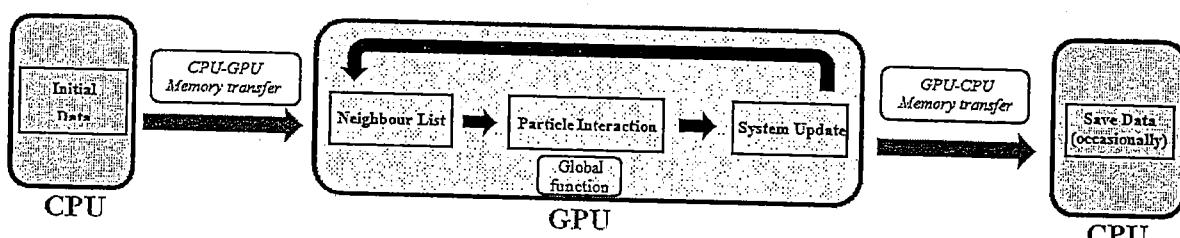


Separate neighbour lists + separate global & device functions):



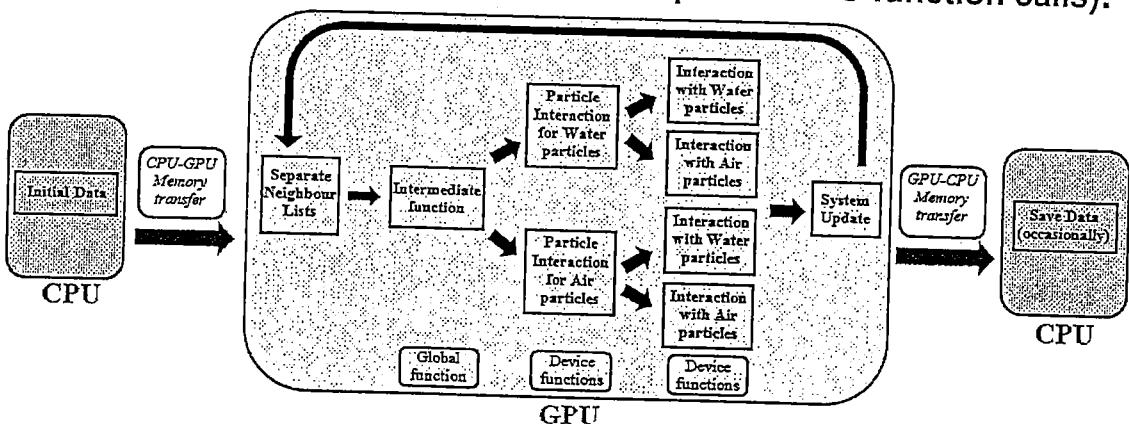
Multi-phase GPU code: separate neighbour lists 2

Original Mono-Phase structure:



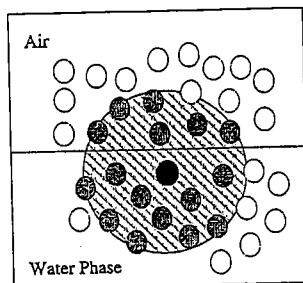
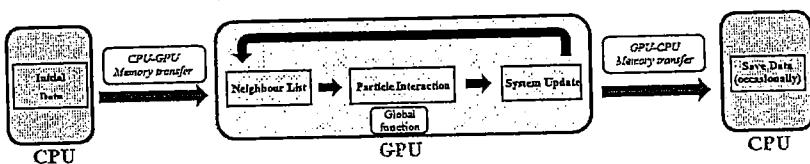
Separate neighbour lists + Intermediate CPU-GPU function

(avoids repeated CPU function calls):



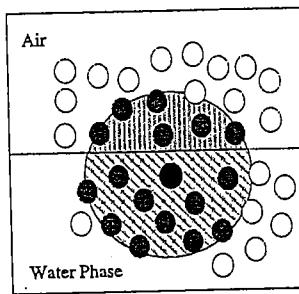
Demands of a multi-phase GPU code

Could we use the original code structure?



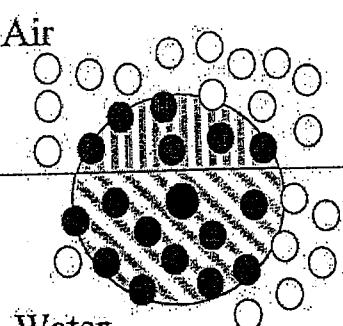
Well, in effect have 4 choices:

1. Use *if*-statements to distinguish each phase
2. Use binary multipliers (0 or 1 for each phase)
3. Separate neighbour lists + separate global & device functions
4. Separate neighbour lists + Intermediate CPU-GPU function
(avoids repeated CPU function calls)



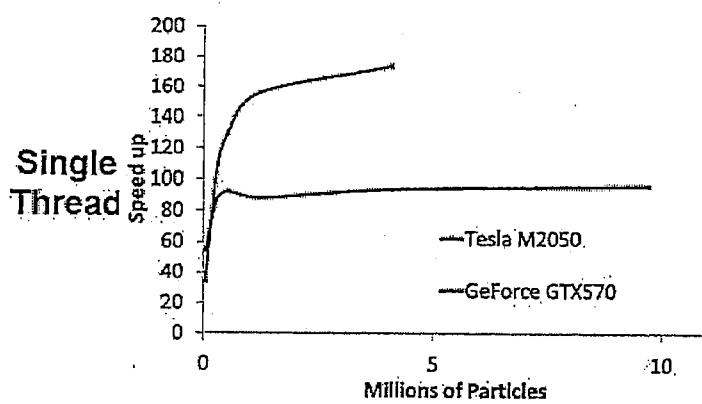
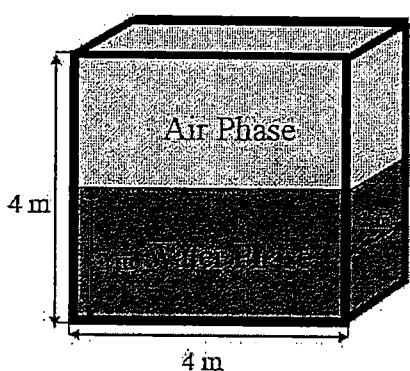
Optimisation of SPH on GPUs

- Calculating inter-particle forces is the most demanding part of SPH
- Research has shown the best practices for optimising:
 - Eliminate conditional *if* statements
 - Reduction of logical operations
 - Minimise CPU-GPU interaction
 - Minimise memory (local and global) transfers
 - Balance computational load on the GPU
- Separate particle and neighbour lists for each phase is beneficial for large particle numbers

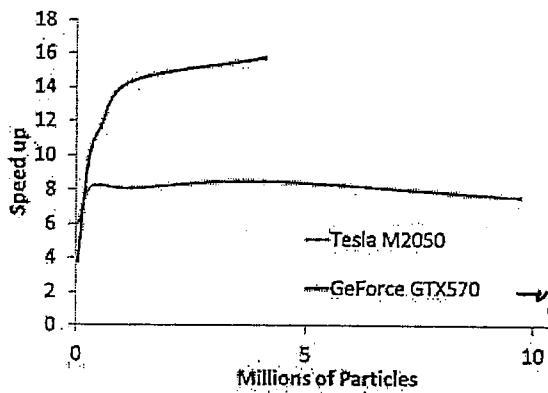


(Mokos et al. 2015)

Runtime Results – 3D

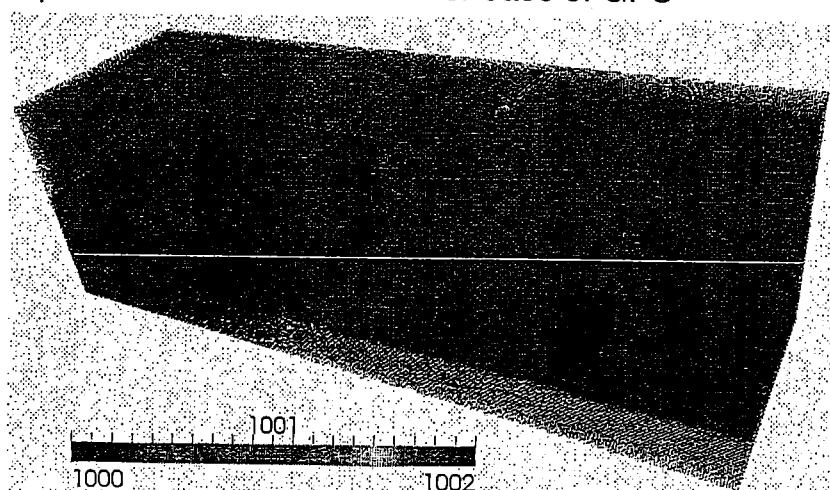


- Results for version 2
- Speed-up is GPU-dependent
- Speedup up to 170 OpenMP compared to single CPU core
- Speedup up to 16 compared to an 8-thread OpenMP computation



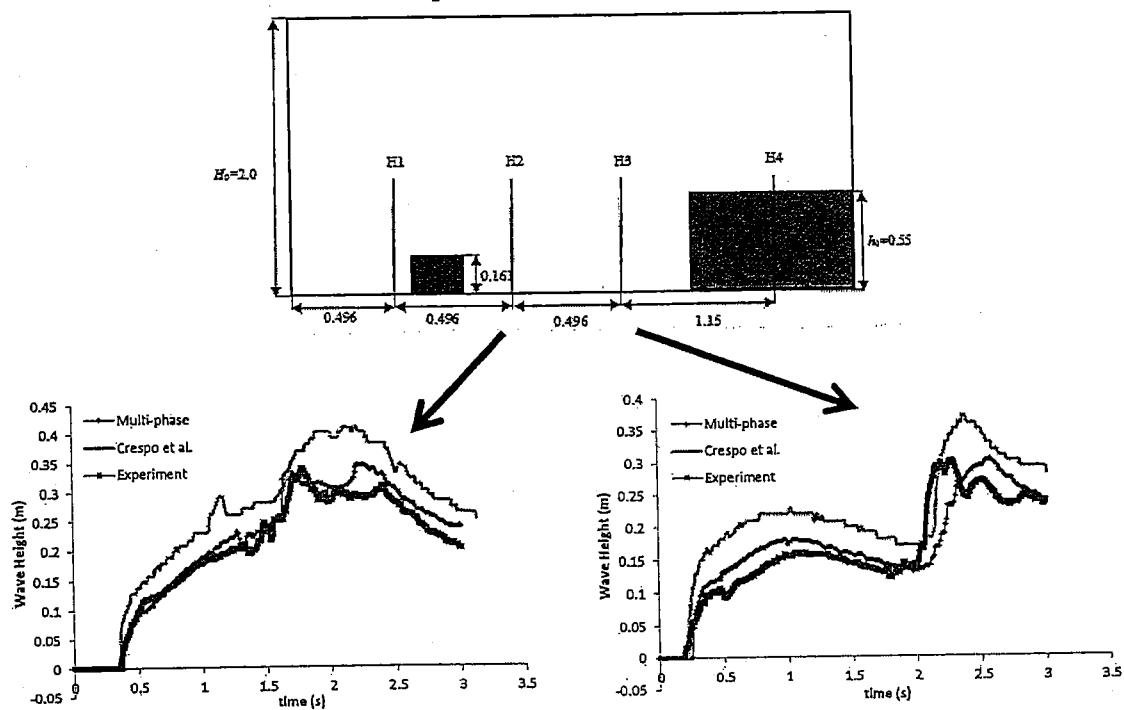
3-D Obstacle Impact: Air & Water (WCSPH)

- Each phase is modelled as weakly compressible.
- Simplest SPH multi-phase formulation chosen for ease of GPU implementation



- Water Density
- Agreement for impact pressures better than monophase
- Simulation runtime for 5 million particles :140 hours for 3s!!!
- (we will come back to this later) Most are AIR particles

3-D Obstacle Impact: Air & Water (WCSPH)



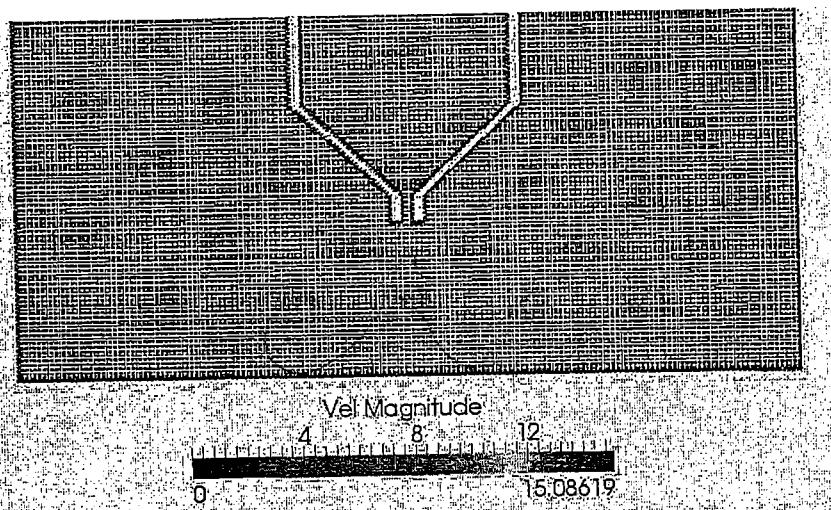
- Agreement for impact pressures better than monophase

Nuclear Applications: mixing

Submerged jet impinging on sediment

- goal: mix/ignite the waste so that it won't explode!

- Configuration
 - $\rho_s = 1.54 \rho_w$
 - $\mu_s^{cr} = 5 \times 10^3 \mu_w$
 - Cohesive sediment
 - 60 000 particles



Multi-GPU for SPH

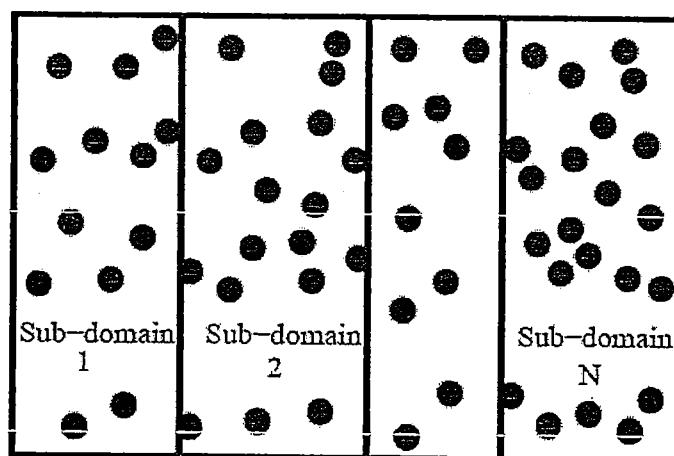
Baldez-Valderas *et al.* (2013)

Dominguez *et al.* (2013)

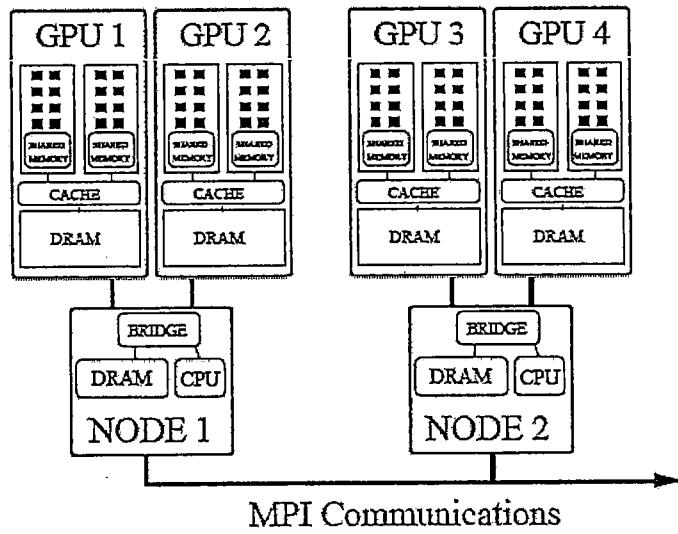


General strategy: spatial decomposition

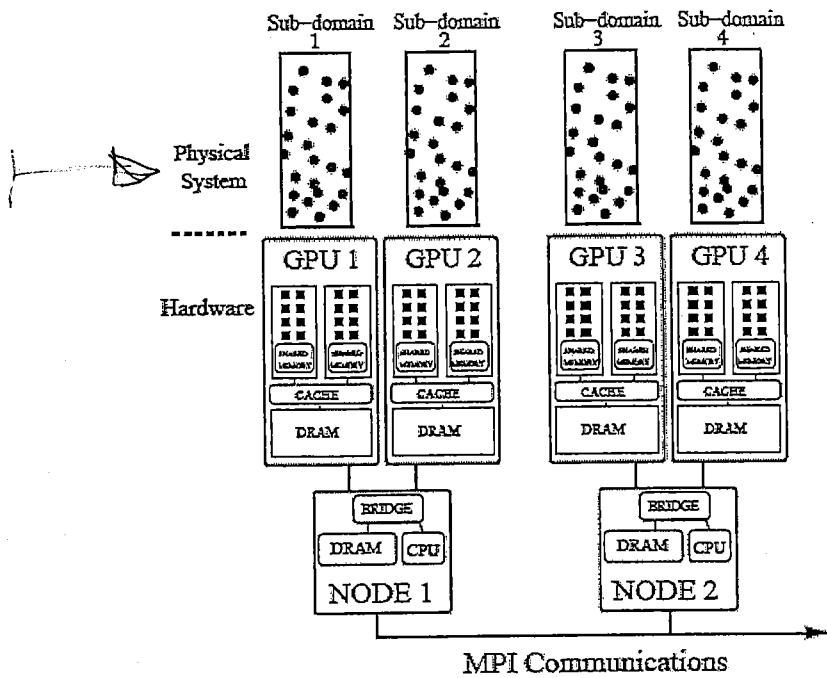
GPU 1 GPU 2 ... GPU N



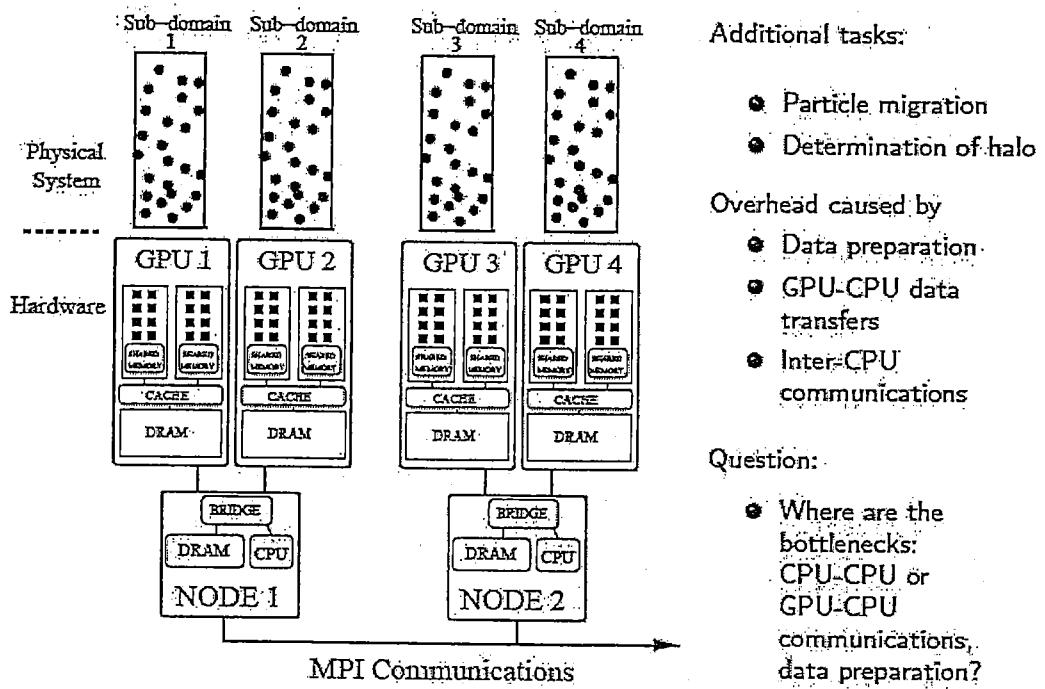
General strategy: a multi-GPU system



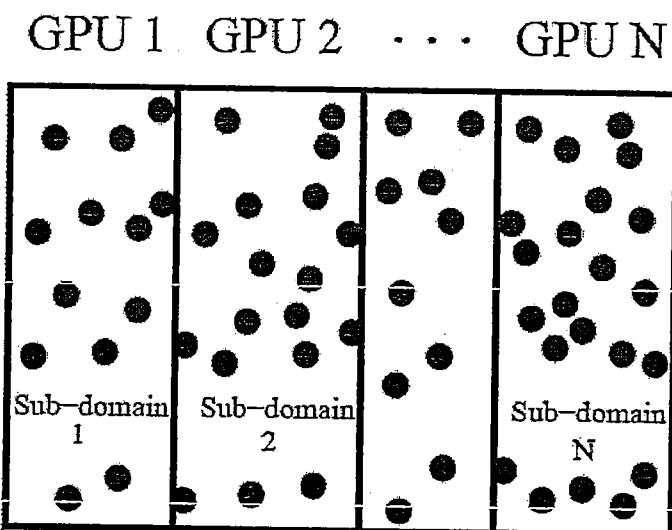
General strategy: a multi-GPU system



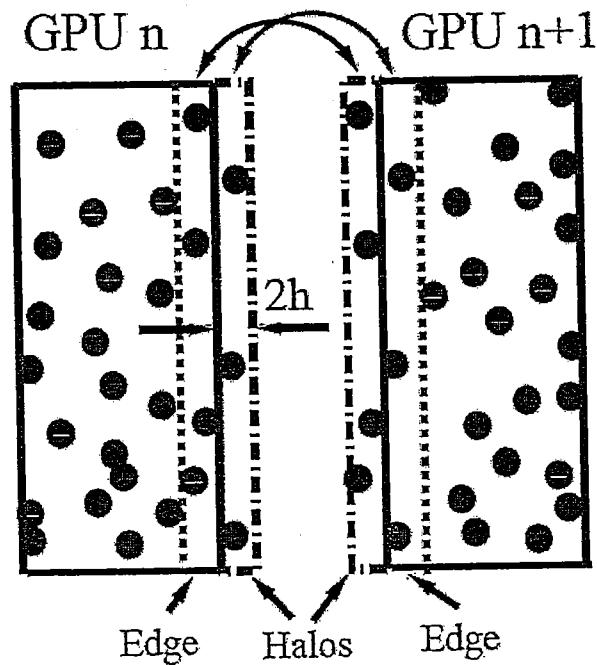
General strategy: a multi-GPU system



General strategy: spatial decomposition



General strategy: halo building

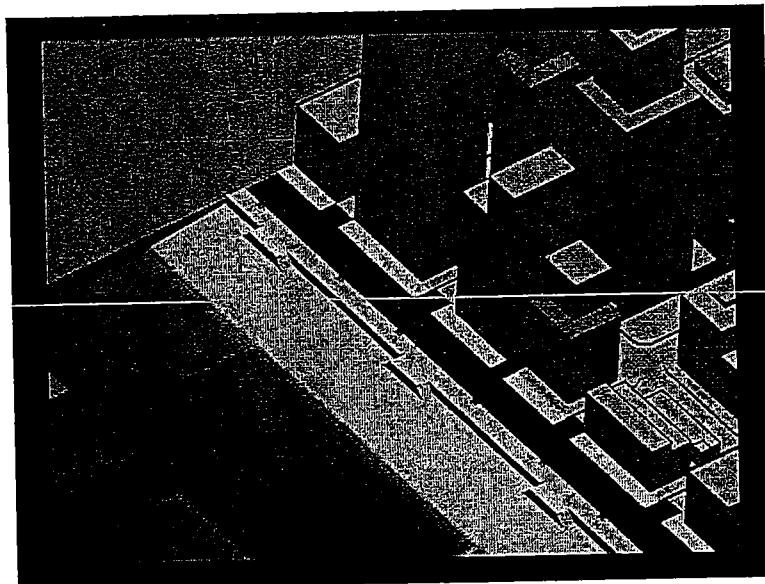


MANCHESTER
IS24

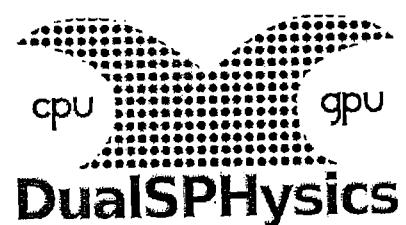
The University
of Manchester

Multi-GPU DualSPHysics

- Dual-SPHysics (cpu & gpu) – SPEEDUPS of 400+ for a Desktop machine = Supercomputer of 1000's cores
- Valdez-Balderas et al. (2013)



GPU4
GPU3
GPU2
GPU1

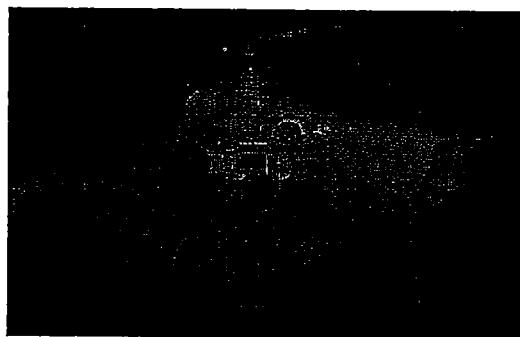
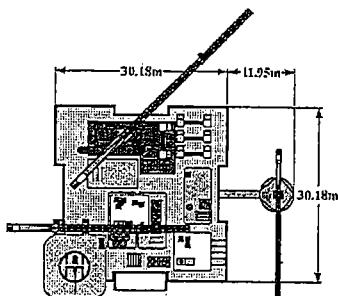


4 GPUs	2.1 hours
1 GPUs	5.7 hours
1 CPU (4 core)	3.5 days
1 CPU (1 core)	2 weeks

5.3 million particles,

SPH Application with MultiGPU DualSPHysics

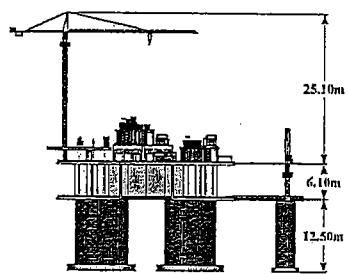
Large wave interaction
with oil rig using 1
BILLION (10^9) particles



[Click here for
Movie 1](#)

Dominguez
et al. (2013)
Computer
Physics
Comm.

We have post-processing tools efficient enough to
visualise the simulation of more than 1000M



$dp = 6 \text{ cm}$
 $np = 1,015,896,172$ particles
 $nf = 1,004,375,142$ fluid particles
physical time = 12 sec
of steps = 237,342 steps
runtime = 91.9 hours

using 64 GPUs
Tesla 2090 of the
BSC

1.5P
64 GPUs

78.15 s
1.5P
1 GPU

Conclusions & Outlook

- SPH is a costly method which needs hardware acceleration
- Like many computational methods, various options exist
- Massive parallelization is possible but really requires sophisticated techniques
- More attractive is to use efficient computational devices
 - → Graphics Processing Units (GPUs)
- Complex physics with Multi-Phase requires careful work

Parallel Shopping with MPI

- SPH has excessively long runtimes: weeks and months
- Uses MPI formalism (MPI = Message Passing Interface)

- à Paris:



Thank you

Acknowledgments

- U-Man: Peter Stansby, Steve Lind, George Fourtakas, Abouzied Nasar
- U-Vigo: Alex Crespo, Jose Dominguez, Moncho Gomez-Gesteira
- U-Parma: Renato Vacondio FHR: Corrado Altomare

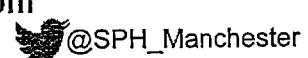
Websites

- Free open-source
- **SPHysics** codes:
<http://www.sphysics.org>
<http://www.dual.sphysics.org>



SPH@Manchester <https://sph-manchester.weebly.com>

International SPH organisation:



SPH research and engineering international community = SPHERIC

<http://spheric-sph.org>

REFERENCES

- Chow, A.D., Rogers, B.D., Stansby, P.K., Lind, S.J., 2018, Incompressible SPH (ISPH) with fast Poisson solver on a GPU, *Computer Physics Communications*, 26, 81-103, doi: [10.1016/j.cpc.2018.01.005](https://doi.org/10.1016/j.cpc.2018.01.005).
- Crespo AJC, Domínguez JM, Rogers BD, Gómez-Gesteira M, Longshaw S, Canelas R, Vacondio R, Barreiro A, García-Feal O. 2015. DualSPHysics: open-source parallel CFD solver on Smoothed Particle Hydrodynamics (SPH). *Computer Physics Communications*, 187: 204-216. doi: [10.1016/j.cpc.2014.10.004](https://doi.org/10.1016/j.cpc.2014.10.004).
- Domínguez, J.M., Crespo, A.J.C., Gómez-Gesteira, M., Marongiu, J.C., 2010, Neighbour lists in Smoothed Particle Hydrodynamics. *International Journal for Numerical Methods in Fluids*. doi: [10.1002/fld.2481](https://doi.org/10.1002/fld.2481).
- Domínguez, J.M., Crespo, A.J.C., Valdez-Balderas, D., Rogers, B.D. and Gómez-Gesteira M. 2013. New multi-GPU implementation for Smoothed Particle Hydrodynamics on heterogeneous clusters. *Computer Physics Communications*, 184: 1848-1860. doi:[10.1016/j.cpc.2013.03.008](https://doi.org/10.1016/j.cpc.2013.03.008).
- Fourtakas G., Rogers B.D. 2016. Modelling multi-phase liquid-sediment scour and resuspension induced by rapid flows using Smoothed Particle Hydrodynamics (SPH) accelerated with a graphics processing unit (GPU). *Advances in Water Resources*, 92: 186-99. doi:[10.1016/j.advwatres.2016.04.009](https://doi.org/10.1016/j.advwatres.2016.04.009).

REFERENCES

- Guo, X., Rogers, B.D., Lind, S.J., Stansby, P.K., 2018, New massively parallel scheme for Incompressible Smoothed Particle Hydrodynamics (ISPH) for highly nonlinear and distorted flow, *Computer Physics Communications*, 233, 16-28, doi: [10.1016/j.cpc.2018.06.006](https://doi.org/10.1016/j.cpc.2018.06.006)
- Longshaw, S.M., Rogers, B.D. 2015. Automotive Fuel Cell Sloshing Under Temporally and Spatially Varying High Acceleration Using GPU Based Smoothed Particle Hydrodynamics (SPH). *Advances in Engineering Software*, 83: 31–44. doi: [10.1016/j.advengsoft.2015.01.008](https://doi.org/10.1016/j.advengsoft.2015.01.008).
- Mokos A, Rogers B.D., Stansby P.K., Domínguez J.M. 2015. Multi-phase SPH modelling of violent hydrodynamics on GPUs. *Computer Physics Communications*, 196: 304-316. doi: [10.1016/j.cpc.2015.06.020](https://doi.org/10.1016/j.cpc.2015.06.020).
- Mokos, A., Rogers, B.D., Stansby, P.K. 2016. A multi-phase particle shifting algorithm for SPH simulations of violent hydrodynamics with a large number of particles. *Journal of Hydraulic Research*, 55 (2), 143-162. doi: [10.1080/00221686.2016.1212944](https://doi.org/10.1080/00221686.2016.1212944).
- Valdez-Balderas D., Domínguez, J.M., Rogers, B.D., Crespo, A.J.C. 2013. Towards accelerating smoothed particle hydrodynamics simulations for free-surface flows on multi-GPU clusters. *Journal of Parallel and Distributed Computing*, 73(11): 1483-1493. doi:[10.1016/j.jpdc.2012.07.010](https://doi.org/10.1016/j.jpdc.2012.07.010).